

FlyDoom: A frozen whole-brain *Drosophila* connectome in a closed sensorimotor loop

Mehmet Utku Öztürk

Independent researcher

mutkuoz@proton.me

September 27, 2026

Abstract

We embed a frozen FlyWire FAFB v783 connectome (139,255 neurons; 2,710,038 edges) in a closed sensorimotor loop with a video game. No parameter is trained; the only fitted quantity is a global synaptic gain calibrated on a non-visual behaviour.

Chemosensory transformations transfer. Sugar drives the proboscis motor pool to 77 Hz, bitter co-stimulation suppresses it by 99%, and odour suppresses a descending neuron one synapse from the lateral horn (-17.7 ± 7.4 Hz), none of it fitted and none surviving a degree-preserving shuffle. Odour also modulates DN_{a02}, the goal-directed steering neuron, by 17.3 ± 1.6 Hz with vision held constant.

Per-cell direction selectivity does not transfer ($|DSI| \leq 0.0121$). Shunting inhibition cancels itself in this neuron model, and making inhibition purely divisive abolishes the residue, so what survives is a threshold conjunction rather than a product. A regional inhibitory scale raises selectivity sevenfold without altering the chemosensory results; a single global value cannot, since it inverts the gustatory response. The point-neuron assumption is itself part of the failure: giving each detector a three-compartment cable, with every input placed along it by its own retinotopic offset, raises the T5 mirror-pair separation an order of magnitude ($0.005 \rightarrow 0.064$). Most of that is electrotonic rather than sequential — shuffled cables retain 0.047–0.050 and a reversed one is better still — and a global excitatory gain of 1.25 reverses the sign, so the model’s preferred direction is not a robust prediction.

Recording site proves more consequential than any parameter. The frozen model carries a 17–26 Hz directional signal at the horizontal system that ablating T4/T5 abolishes, and the regional scale that maximises per-cell selectivity drives it to zero. That signal reaches DN_{p15}, the optomotor descending neuron, at up to 32 Hz, while the steering readout used throughout was DN_{a02}, which draws 2.3% of its input from vision and fires near its refractory ceiling.

Decoding yaw from DN_{p15} with every correction applied, the model’s commands beat a command-matched random agent on health collected ($+12.67 \pm 10.80$), survival, exploration and stuck-fraction over thirty held-out seeds. Mirroring the retina abolishes that advantage, and so does freezing it. The advantage is vision-dependent, though the model’s own decline without vision is not individually significant and no recognisable reflex underlies it: an optomotor drum, a looming stimulus and a fixation target each produce no directional response. In a purpose-built bar arena the model does hold a dark bar at a fixed retinal azimuth, but a frozen retina holds it as tightly and an empty cylinder scores the same; swept open loop, the turning command’s dependence on bar position is smaller than what the empty cylinder and the frozen retina produce. Against a size-matched static disc, the looming response is 26.42 against 26.44 Hz at LC4: the expansion adds nothing measurable. Three errors in the interface — an eye map read with the wrong hexagonal convention, a reversed steering sign, and a walk command that ran the agent backwards — partly produced that advantage. Corrected, over 120 seeds, it survives only in an arena rebuilt so the fly can resolve its structure ($+5.67 \pm 4.88$ health, abolished by mirroring, weakened by freezing) and vanishes where the walls are finer than its acceptance angle; a second block of sixty seeds halves the effect the first reported, so one block of sixty pinned it no better than to a factor of two; the corrected eye doubles the population direction signal and still produces no optomotor response.

One sensory addition does reach behaviour. Airflow, delivered as physics — a source’s odour reaching only a fly downwind of it, and the flow direction arriving at the Johnston’s organ afferents already present in the connectome — moves DN_{p15}’s left–right differential by 8 Hz with consistent sign across seeds, and reduces the model’s net rotational bias by a quarter while straightening its path by a fifth. It does not orient the animal to the wind (resultant 0.639 with airflow against 0.549 without)

and does not help it find food. It stabilises course, which is the function optic flow serves in a fly and which this model’s motion pathway cannot provide.

Three results generalise. Any stimulus manipulation is an invalid closed-loop control, as degree-preserving shuffles already are, because an arm that behaves differently samples a different stimulus; command-matched controls and trajectory-matched replay avoid this. Episode-level measures failed to resolve a population signal varying tenfold across conditions. And a square-wave stimulus places its phases half a period apart on average, so equal sample counts do not remove drift from a phase-locked comparison.

Keywords: connectomics; *Drosophila*; whole-brain simulation; closed-loop embodiment; elementary motion detection; olfactory valence; negative results

1 Introduction

An electron-microscopy reconstruction of an adult *Drosophila* brain now provides every chemical synapse between all 139,255 neurons [Dorkenwald et al., 2024], together with cell-type annotation [Schlegel et al., 2024] and per-synapse neurotransmitter prediction [Eckstein et al., 2024]. This raises a question that was previously unaskable: *how much of an animal’s behaviour is implied by its wiring alone?*

Two recent lines of work answer parts of it, and they answer differently. Shiu et al. [2024] built a whole-brain leaky integrate-and-fire model with a uniform parameter set and a single fitted scalar, the per-synapse weight W_{syn} , and recovered gustatory behaviours such as proboscis extension. Lappalainen et al. [2024] built a connectome-*constrained* model of 64 optic-lobe cell types in which the unknown biophysical parameters were *optimized* by gradient descent on a motion-estimation task, and recovered single-neuron responses matching 26 experimental studies. These differ in two ways at once: whole-brain versus optic lobe, and frozen versus optimized.

Both are driven open loop. Wang-Chen et al. [2024] close the loop, embedding the optimized visual network of Lappalainen et al. [2024] in a biomechanical fly with reinforcement-learned controllers. The configuration that has not been tested is the remaining corner: a *whole-brain, frozen, unoptimized* connectome placed in a closed sensorimotor loop.

That corner is informative precisely because of the contrast. If a frozen whole-brain model reproduces a computation, the wiring plus a global gain sufficed. If it fails at a computation that the optimized model achieves on the same wiring, the difference isolates what optimization supplied.

We report such a test. We drive the frozen connectome with the rendered frames of *Doom* [Kempka et al., 2016], read eight characterized descending neurons, and return their activity to the game as control input, in real time. The environment is adversarial and non-stationary but, crucially for interpretation, entirely outside the fly’s evolutionary experience, so nothing here is a test of ecological competence. It is a test of which sensorimotor transformations survive the transfer from wiring diagram to running system.

Contributions.

1. A closed-loop embodiment of a frozen whole-brain connectome (139,255 neurons, single consumer GPU or CPU), released as open code, running at $0.19\times$ realtime.
2. A modality dissociation: chemosensory transformations transfer, visual motion computations do not.
3. A mechanistic account of the visual failure. Twenty candidate explanations were implemented and measured; the operative one is that shunting inhibition cancels itself, so the neuron sums its inputs where the computation requires a product. Separating the two components of inhibition shows the residue is a threshold conjunction rather than a division, and that the measured reversal potential is already at the optimum.
4. A demonstration that the standard single-cell measure of direction selectivity and the population signal at the optomotor output stage move against each other under the same parameter, so that




Figure 1. The closed loop. Rendered frames are sampled by the real ommatidial lattice, propagated through the frozen connectome for 57 leaky integrate-and-fire substeps per game tic, and read out at eight characterized descending neurons whose activity becomes the next action. Taste and smell enter as direct drive on gustatory and olfactory receptor neurons. No component is trained and no parameter is fitted to the environment.

a search maximising the former removes the latter. With the recording site corrected, the frozen model carries a directional signal that survives ablation of the motion detectors and that the reported parameterisation destroys.

5. A characterisation of the output stage, which no parameter in the sensory pathway can compensate for: four fifths of descending neurons never fire, the steering readout runs near its refractory ceiling on 2.3% visual input, and the directional signal arrives instead at the horizontal system’s actual descending target.
6. A parameterisation carrying no fitted quantity. Published per-receptor conductances restricted to the visual populations preserve both chemosensory results and recover most of the selectivity gain, where the same values applied brain-wide silence a descending neuron outright, which isolates scope from the values themselves. Across every parameter that does not entail per-cell-type fitting, selectivity does not exceed 0.0792, so the shortfall is a measured property of the accessible space.
7. A demonstration that a single global synaptic gain is inconsistent rather than merely approximate. Raising the inhibitory-to-excitatory ratio restores the multiplicative interaction, but the optic lobe requires a value the subesophageal zone does not tolerate; applied regionally the same factor increases direction selectivity sevenfold at no cost to the chemosensory results. The same manipulation does not produce closed-loop behaviour, which we report as a negative result: no parameterisation separates from a command-matched control by more than the frozen model does.
8. A positive result that beats a degree-preserving shuffle: an olfactory channel reaches descending neurons through the anatomically named route with the wiring-predicted sign.
9. A methodological caution: shuffle controls are invalid in closed loop.

2 Results

2.1 Model and closed-loop configuration

We load FlyWire FAFB v783 and retain edges of at least five synapses, giving 2,710,038 edges carrying 31,578,726 synapses among 134,013 connected neurons (Table 1). Signs follow the predicted transmitter, with glutamate treated as *inhibitory*: it acts through GluCl chloride channels in flies [Cleland, 1996] and

Table 1. Graph as simulated. Unthresholded synapse count matches the figure published for this release, confirming the extraction.

Neurons (annotated)	139,255
Edges, unthresholded	16,847,997
Synapses, unthresholded	54,492,922
Edges simulated (≥ 5 synapses)	2,710,038
Synapses simulated	31,578,726
Neurons with ≥ 1 edge	134,013
Graded (non-spiking) units	77,873
Resident GPU memory	31 MB
GPU kernel time per 0.5 ms step	$\sim 310 \mu\text{s}$
Wall-clock per 0.5 ms step	1.4 ms (CPU, 8 threads, CSR delivery)
Closed-loop rate	$0.19 \times \text{realtime}$

accounts for 16.6% of the edges we simulate (19.5% before thresholding), so importing the vertebrate convention would mis-sign about a sixth of the brain.

We sign per edge rather than per neuron, and in this pipeline the choice is not free. Transmitter is predicted per synapse and reported per neuropil, so one neuron can carry different predictions in different neuropils: 42,078 of the 128,972 presynaptic neurons in the simulated graph (32.6%) carry more than one, and 26,781 (20.8%) carry both signs. Collapsing each neuron to a synapse-weighted majority, the usual way of imposing Dale’s law, would flip 3.1% of edges and 2.6% of synapses. This matters for reading the literature as much as for the model: the alternative Princeton release has Dale’s law applied upstream (0 of its 139,003 presynaptic neurons carry more than one transmitter), so per-neuron and per-edge signing coincide there and not here, and transmitter statistics quoted from one release do not describe the other.

Photoreceptors required an explicit override. They release histamine [Hardie, 1989], which is not among the six transmitters the classifier predicts; on R1-6 output edges the classifier returns ACh 71.2%, Glu 20.2%, GABA 5.8%, and every excitatory call yields the wrong sign for a histaminergic chloride synapse. We force all photoreceptor output inhibitory, flipping 11,920 of 18,437 edges. Without this the optic lobe processes a contrast-inverted image.

Neurons are leaky integrate-and-fire with the parameters of Shiu et al. [2024], themselves drawn from Kakaria and de Bivort [2017]: $V_{\text{rest}} = -52 \text{ mV}$, $V_{\text{thresh}} = -45 \text{ mV}$, $\tau_{\text{mem}} = 20 \text{ ms}$, $\tau_{\text{syn}} = 5 \text{ ms}$, $t_{\text{refrac}} = 2.2 \text{ ms}$, $t_{\text{delay}} = 1.8 \text{ ms}$. We integrate the linear subthreshold system exactly rather than by forward Euler, which we measured 6% low on a single postsynaptic potential. At $dt = 0.5 \text{ ms}$ this gives 57 substeps per game tic, with the frame held across all 57.

Vision enters through the real retinotopy: `column_assignment` supplies 796 ommatidial columns per eye on a hexagonal lattice. Each column samples the rendered frame through a Gaussian acceptance function, implemented as a separable pre-blur followed by bilinear sampling at that column’s gaze direction. The engine draws a planar perspective projection, so a column at azimuth θ is sampled at screen coordinate $\tan \theta / \tan(\text{FOV}/2)$ rather than linearly in θ ; the linear map misplaces a column at 10° of gaze by 87% and one at 30° by 71%, converging only at the very edge of the viewport. Since retinotopy is the substrate every motion computation runs on, and a warp of that size makes the angular spacing between neighbouring columns vary about twofold across the field, we note it as a correctness requirement rather than a refinement: rigid motion of the world then sweeps the retina at a speed that depends on where one looks, which is precisely the input a correlator tuned to a fixed spacing cannot use. The two eyes are splayed $\pm 40^\circ$; collapsing them onto the screen centre makes both eyes see the same image and drives the left–right steering difference to zero by construction. The engine’s field of view must be set through the console rather than the command line, passing `+fov` is silently ignored, and renders at 90° , 130° and 160° are bit-identical, and it resets with each episode. At 130° the viewport covers 64% of the lattice and at 160° , 80%; remaining columns are filled with the frame’s mean luminance, because a dark surround is a permanent high-contrast edge that the looming detectors report as a stationary object. The full-eye runs at the end of Sec. 2.13 remove this limit, so that

every column samples the scene (Methods).

Frames are not held across the substeps of a tic. Doom draws at 35 Hz and the simulator runs 57 substeps per frame, so a held frame presents the retina with a step change followed by 56 substeps of exactly zero temporal derivative, while the delayed arm of the correlator is 80 ms, or three frames: that is, an impulse train aliased against the detector’s own delay line. Because a fly in a room receives continuous motion and the strobe is an artifact of the game engine rather than of the biology, we interpolate luminance linearly between consecutive frames across the substeps. This also makes the retina’s Weber adaptation advance once per substep as its time constant assumes; stepping it once per tic instead gave it an effective τ of 14 s in place of 0.25 s.

Output is read from eight characterized descending neurons, DNa02 for steering [Rayshubskiy et al., 2025], BPN for forward walking [Bidaye et al., 2020], MDN for backward walking [Bidaye et al., 2014], DNp01 (the giant fibre) for escape [von Reyn et al., 2014], plus the proboscis motor pool and gustatory receptor neurons. The graded channels use the engine’s delta buttons; binary actions use a two-threshold Schmitt trigger.

Two of the graded channels read a *difference between the two sides* of a bilaterally paired cell type, and there the standing asymmetry of a single reconstructed brain is a confound rather than a command: DNa02 sits at 150 against 257 Hz and DNp01 at 110 against 206 Hz, both of which exceed their command clamps and pin the button. We therefore subtract a slow baseline ($\tau = 3$ s) from the paired channels before decoding, keeping the transients and discarding the DC, and set each gain so that three standard deviations of what remains reach the clamp. The forward channel is not a bilateral pair, it is BPN – MDN, two opposing cell types, so its standing +27.6 Hz is a tonic walk command and is kept, with the gain placing it at half the clamp instead. We note the consequence rather than manage it: that difference carries ± 1.3 Hz of modulation on a 27.6 Hz mean, so walking speed here is 4.6% modulated and effectively tonic.

2.2 Gustatory pathway

As a positive control we reproduce the gustatory result of Shiu et al. [2024]. Stimulating sugar receptor neurons at 100 Hz drives the proboscis motor pool to 77 Hz; stimulating bitter receptors drives it to zero; stimulating both together yields 99% suppression relative to sugar alone.

This control has a known circularity that we state rather than manage: like Shiu et al. [2024] we fit W_{syn} on this very behaviour, obtaining 0.1655 mV against their 0.275 mV (the difference is attributable to a different connectome release and synapse pipeline). What the test genuinely establishes is therefore weaker than it appears, that *some* single scalar gain reproduces the behaviour, but it is not vacuous: the bitter–sugar suppression ratio was not fitted, and a sign-scrambled graph admits no such scalar.

2.3 Olfactory valence

The environment renders light and nothing else, but a fly in a room with a large animal in it smells that animal; simulating vision alone makes the model impoverished rather than conservative. We therefore added a channel feeding two real labelled lines, ORN_DA1, carrying the pheromone cVA [Kurtovic et al., 2007], and ORN_DM1, carrying vinegar [Sammelhack and Wang, 2009], with enemies mapped to the former and health pickups to the latter.

The channel is deliberately built weak so that it cannot substitute for vision. Most importantly it carries *no direction*: a fly’s antennae are ~ 0.3 mm apart and cannot triangulate, so both sides receive bit-identical drive. It is additionally saturating, slow ($\tau = 0.2$ s), patchy (3 Hz bursts at 60% duty), and persistent ($\tau = 1.5$ s), so a source moving out of sight still smells.

Testing this in closed loop would be invalid (Section 2.4), so we ran it *open* loop: the agent is held still and the game stepped with a null action, so enemies still move and the scene still evolves, but it evolves identically in both arms because the seed is fixed and nothing feeds back. Smell is then the only difference between arms, and LC4, a visual cell that should not move, serves as a built-in check that the comparison is sound.

The lateral horn, the innate-valence pathway [Jefferis et al., 2007, Dolan et al., 2019], is dormant at 0.32 Hz without smell and rises to 14.1 ± 5.5 Hz with it, in all ten seeds (Table 2). The effect reaches the descending neurons in *one synapse*, and its sign is predicted by the wiring: LH $\rightarrow$ DNp01 is –248

Table 2. Effect of the olfactory channel, open loop, 10 seeds $\times$ 320 tics. Δ is the mean on-minus-off difference in Hz; “sign” counts seeds agreeing with the mean direction, with p from a two-sided sign test. In the intact connectome every motor and valence population moves consistently; under a degree-preserving shuffle of the same graph every effect falls to the 0.00–0.11 Hz scale of the LC4 control. Sign consistency alone does not separate the arms — the shuffle retains some of it — so the comparison to read is magnitude. LC4 is a visual cell whose activity should be unaffected by smell; it carries a small offset in both arms, at $10479\times$ smaller magnitude than the lateral horn response it bounds.

Population	Intact				Degree-preserving shuffle			Role
	Δ (Hz)	s.d.	sign	p	Δ (Hz)	s.d.	sign	
LHN	+14.06	5.46	10/10	0.002	+0.01	0.01	8/10	innate valence
DNp01	−17.67	7.42	9/10	0.021	−0.00	0.27	6/10	escape (giant fibre)
DNa02	+11.68	5.29	9/10	0.021	+0.11	0.34	7/10	steering
BPN	+1.21	0.46	10/10	0.002	+0.07	0.04	10/10	forward walking
MDN	−0.96	0.40	10/10	0.002	−0.02	0.15	6/10	backward walking
LC4	−0.00	0.02	8/10	0.109	+0.07	0.03	10/10	<i>visual control</i>

synapses, hence inhibitory, and DNp01 is suppressed by 17.7 ± 7.4 Hz, in nine seeds of ten. Meanwhile LC4, a visual population that should not move at all, does not: it shifts by -0.00 ± 0.02 Hz, four orders of magnitude below the lateral horn response it bounds.

Crucially, a degree-preserving shuffle of the same graph, run through the identical protocol and the identical seeds, does not reproduce this (Table 2, right). The valence and steering effects collapse by two to three orders of magnitude, LHN from +14.06 to +0.01 Hz ($1563\times$), DNa02 $102\times$, DNp01 to within noise of zero, landing at the 0.00–0.11 Hz scale of the LC4 control, that is, at the network’s generic coupling floor.

Sign consistency and effect magnitude discriminate differently here, and only the second separates the arms. Sign consistency does *not*: several shuffled populations remain fairly sign-consistent across seeds, including LC4 at 10/10, the one population that should not be responding at all. What discriminates is magnitude. A sign test on a 0.01 Hz effect is a statement about a systematic that is three orders of magnitude too small to matter, and we interpret it accordingly. The shuffled DNp01 residual reversed sign between two independent executions of the protocol, where every intact effect held its direction, the behaviour expected of noise rather than of a transformation. In summary: the intact effects are large and reproducible while the shuffled ones sit at the control’s floor. *The effect requires the wiring.*

The BPN and MDN walking effects support a weaker claim than the rest. They are consistent in direction across all ten seeds, but at ± 1 Hz they are only $\sim 17\times$ above the shuffled floor rather than $\sim 100\times$, and we draw no conclusion from them in isolation.

The shuffle also changes the baseline in an informative way. Without smell the intact lateral horn sits at 0.32 Hz; the shuffled one idles at 25.7–25.8 Hz across seeds, an $80\times$ higher floor. Preserving degree while destroying specificity does not merely remove the response: it removes the *quiet*. The intact lateral horn is selectively silent and selectively driven, which is the behaviour a labelled-line valence pathway should show, and neither property survives rewiring.

Two caveats apply to this result throughout. Odour source distances are read from the engine’s label buffer, so the model was *told* that enemies exist; this is evidence about what the connectome does with a signal, not evidence that it can detect one. And 1,227 lateral horn neurons pooled together mixes many functional channels, so “the lateral horn responds” is coarser than it sounds.

Synaptic weight against path length

We first routed the same odour at the pC1 aggression neurons, on the strength of a three-hop path touching 8 of 10 cells, and it failed flat: 0.00 Hz across a tenfold sweep of drive. Measured afterwards: the reason is structural: DA1’s contribution is a rounding error against pC1’s balanced +530/−555 of input, and pC1 is a bistable latch that must be pushed over rather than nudged. The lateral horn is the




Figure 2. Olfactory valence reaches the descending neurons, and needs the wiring to do it. **(a)** Per-seed lateral horn rate with smell off and on in the intact connectome; all ten seeds rise. **(b)** The same protocol on a degree-preserving shuffle of the same graph: no seed moves appreciably, and the shuffled baseline idles near 21 Hz where the intact pathway is dormant at 0.32 Hz. **(c)** Mean change per population, both arms (error bars are s.d. over ten seeds). LC4 is a visual cell included as a control: it should not move, and does not.

opposite on every count: it receives 21% of the odour relay’s entire output, has no latch, and sits one hop from the output. Hop count is a poor proxy for influence in a connectome; synaptic weight against competing input is the quantity that matters.

2.4 Validity of shuffle controls in closed loop

Degree-preserving shuffles are the standard control for attributing an effect to connectivity, and we use one above. But in a closed loop they silently stop being a control.

We measured this directly. Over one matched episode the intact agent encountered an enemy on 186 of 500 tics and finished at full health; the shuffled agent encountered one on 489 of 492 tics and finished at 16 health. The point this episode makes is qualitative, that the arms visit different states, and it does not require the difference to be of any particular size; but see the power note below before reading episode counts as effects. The two arms behaved differently, and because behaviour determines what the agent sees, they were no longer sampling the same stimulus distribution. Any cross-arm correlation then compares two different worlds rather than two different brains.

The consequence generalizes to any embodied connectome study: *controls must hold the stimulus fixed*, which in practice means running the comparison open loop, as we do for the olfaction result.

2.5 Motion vision and looming

Direction selectivity fails. Presenting drifting gratings and measuring the direction selectivity index $DSI = (R_{\text{pref}} - R_{\text{null}}) / (R_{\text{pref}} + R_{\text{null}})$ in T4/T5, the canonical elementary motion detectors [Maisak et al., 2013], gives $|DSI| \leq 0.0121$. Experimental studies routinely *select* terminals for analysis at $DSI > 0.5$, so the model sits some $40\times$ below the threshold at which a real terminal would merely qualify as direction-tuned.

Looming fails in the same way and more starkly (Table 3). Neither the intact nor a degree-preserving shuffled connectome distinguishes an expanding from a contracting stimulus. The comparison that shows this is against a *size-matched* static disc, held at the sweep’s mean angular size so that the two stimuli deliver the same total drive and differ only in that one expands. Against it, LC4 gives 26.42 Hz to the looming sweep and 26.44 Hz to the static one: to four significant figures the expansion adds nothing, and what the pathway reports is mean angular size. (A static disc held at the sweep’s *final* size, the control used in earlier versions of this table, delivers far more drive than the stimulus it controls for and can therefore beat looming however well the detectors work; we report it only for continuity.)

Table 3. Looming (M4), open loop, mean rate in Hz; $l/|v| = 40$ ms, 1 s, driving L1+L2+L3. *Static (matched)* holds the sweep’s mean angular size and is the control the comparison needs; *static (final)* holds its end point and is the defective one. Against the matched control the expansion adds nothing, in either arm, and a blank field reproduces most of the response.

	Intact			Shuffled		
	LC4	LPLC2	DNp01	LC4	LPLC2	DNp01
Looming	26.42	0.00	131.5	5.38	2.27	71.5
Static (matched)	26.44	0.00	141.0	5.16	2.10	70.0
Receding	24.31	0.00	147.0	4.16	1.79	65.5
Blank	24.90	0.00	120.0	4.26	1.64	69.5
Static (final)	77.21	0.00	145.5	47.15	15.30	213.0

The blank field is 94% of the looming response, and the shuffled graph behaves the same way, so the measurable looming response is not attributable to connectivity at all.

DNp01, the giant fibre, fires at 120–147 Hz in every condition including the blank, which is not a looming response that is too weak but a cell that is saturated: no behaviour downstream of it in this model carries information about the stimulus.

The same measurement answers a question about the model’s locomotion. The forward command is BPN – MDN, and MDN — the moonwalker descending neuron, which commands backward walking [Bidaye et al., 2014] — is read as a reverse command by the closed-loop decoder, so what drives it is a behavioural question and not only a question about the escape circuit. It is not driven by looming: MDN gives 3.50 Hz to a blank field, 4.50 to the size-matched static disc, 5.25 to the looming sweep and 6.25 to the large static disc, which is the rank order of mean dark *area* and the same order as LC4. Nor is it driven by contact: under forced antennal stimulation (three seeds, 200 tics each) it moves from 7.18 Hz with no contact to 6.50 on the left and 8.26 on the right, about a hertz and not consistently in one direction, with BPN flat at 30.5–30.7. In *Drosophila* backward walking is evoked both by looming and by the antennae. In this model it is evoked by neither, and the agent reverses when a wall fills its eye.

State of the looming detector

In closed loop LPLC2 reads exactly 0.00 Hz, and the reason is measurable. It sits at -52.2 mV, below its own resting potential and 7.2 mV from threshold, because its heaviest input, PVLP011 at -45.6 synapses per cell, fires at 371 Hz. The model’s absolute refractory ceiling is $1/t_{\text{refrac}} = 455$ Hz, so that one inhibitory cell runs at 82% of the maximum rate the model permits, and the inhibitory conductance it delivers exceeds the excitatory arm (Tm5f, +44.7 synapses per cell but only 36 Hz) by roughly tenfold. This is the T4a arm imbalance again, in a second and independent circuit.

That it is an imbalance rather than a shortage is what makes it diagnostic. We swept every input-side scale available (photoreceptor drive, the graded rate ceiling, and Weber contrast gain) and LPLC2 does not leave 0.00 Hz at any setting. Reducing drive lowers PVLP011 from 371 to 113 Hz but takes LC4 from 53 to 0.0 Hz with it. The result is expected on inspection: an input scale multiplies both arms of a ratio and cancels, so no common gain can change a 10:1 imbalance. Only a gain that acts differentially on the two arms can, which is to say a per-type gain: the quantity the connectome does not specify and the optimized model of Lappalainen et al. [2024] obtains by fitting.

A tonic depolarising drive to the optic lobe does release it, and shows the same trade rather than escaping it. Raising the bias from 0 to 6 mV lifts LPLC2 from 0.0 to 64.7 Hz, but it also lifts MDN, and the walking command is BPN – MDN: that difference falls from +22 Hz to +0.3 Hz by 2 mV and reverses sign by 6 mV, at which point the model walks backwards. Meanwhile the population above 200 Hz grows from 197 cells to 991. There is no setting at which the looming pathway conducts and locomotion survives.

A blank-field control shows the trade is worse than that, and it is the reason we report bias sweeps

with one. Presenting an *empty screen* at 7.5 mV bias drives LPLC2 to 76.17 Hz, against 76.22 Hz for an expanding disc, a stimulus contribution of 0.07%. At 4 mV it is 30.52 blank against 30.62 looming. The bias that releases the cell is the same quantity that determines its output: past threshold we are no longer measuring vision but measuring the bias, and the fraction of the response attributable to the picture *falls* as the bias rises. There are two regimes and neither is looming detection: LPLC2 silent, or LPLC2 firing at a rate the stimulus barely perturbs. Reporting the looming response alone, without the blank, would have made 76 Hz look like success.

We also tested whether the failure is an artefact of where visual drive enters. By default we inject at the lamina monopolars, one synapse past the photoreceptors, which bypasses two real features of the eye: neural superposition, in which six photoreceptors sharing a visual direction converge on one cartridge, and the histaminergic first synapse that inverts the signal. Injecting at R1-6 instead restores both. It changes nothing. At matched bias the two sites agree to within the blank-field offset (LPLC2 18.4 against 30.6 at 4 mV, 76.2 against 76.4 at 7.5 mV), and at zero bias photoreceptor injection is silent throughout, because the inhibition-dominated optic lobe has no baseline to disinhibit from. The one thing photoreceptor injection does carry is a bias of its own: FAFB's uneven optic-lobe proofreading recovers a column for 451 of 785 left columns against 749 of 796 right, so it is unusable for any measurement of a left-right difference, which is why the default sits at the lamina.

Across every condition in that sweep, looming and receding differ by less than the blank-field offset, and the *sign* of the difference flips between conditions. The only regime with genuine stimulus-driven structure is zero bias at the lamina, where LC4 gives 77.2 Hz for a static disc against 24.9 Hz blank, a real threefold response, to dark *area*. The model reports how much dark is present, never whether it is growing.

Downstream visual functions fail correspondingly. LPLC2 activity is uncorrelated with true angular size ($r = +0.000$), and no azimuth signal exists anywhere we could find: neither LC11 ($r = +0.034$) nor LC4 ($r = +0.070$) carries target bearing. LC10a, the cell type that would carry it, is silent at 0.00 Hz, receiving only +45 / - 83 synapses per cell against LC11's +477 / - 456 and LC4's +439 / - 259; it is held down by A0TU042 and TuTuAa. Its known state-dependent gain boost is a *male* courtship mechanism gated by P1, and FAFB is a female brain, so that circuit is absent. Nor is there a state signal to substitute for it. A0TU042 suppresses LC10a at -18 synapses per cell, but pC1 and A0TU042 are not connected in this brain at all, zero edges in either direction, at any synapse threshold, so the aggression state cannot reach that suppression to lift it or to deepen it.

2.6 Candidate mechanisms for the motion failure

Direction selectivity and looming selectivity are both computations over the *temporal order* in which neighbouring retinotopic columns activate. Motion to the right makes column *A* lead column *B*; motion to the left makes *B* lead *A*. The two cases contain identical events and differ only in order.

A symmetry argument appears to make the failure structural. Presynaptic inputs sum into a conductance and the only nonlinearity is a threshold on that sum; since $\text{thr}(A + B)$ is symmetric in its arguments, no unit of this class would distinguish the two cases. The argument is invalid, because it holds only for a unit with no memory. The model has memory in two places the argument ignores: per-connection conduction delays, and a synaptic conductance that decays with $\tau_{\text{syn}} = 5$ ms rather than vanishing within the step. A delayed arm and a persistent conductance make the summand at any instant depend on stimulus history, and the symmetry does not hold.

We settled this by construction rather than by argument. Using the same neuron model, the same integrator and the same time step, we built a two-input correlator (two input units, one delayed, converging on a third) and measured it under the identical grating protocol. It reaches $\text{DSI} = -0.79 \pm 0.017$ (30 s $\times$ 5 seeds), against 0.002–0.016 for T4a in the network. The model class can compute direction selectivity to a magnitude far above the experimental selection threshold. Whatever prevents the network from doing so, it is not the arithmetic of summation and threshold. Biological elementary motion detectors combine a delayed and an undelayed arm *multiplicatively* [Hassenstein and Reichardt, 1956, Borst et al., 2010], realized in *Drosophila* as preferred-direction enhancement together with null-direction suppression [Haag et al., 2016, Gruntman et al., 2018]; a delay plus a saturating threshold approximates that product well enough to yield the 0.79 above.




Figure 3. The motion failure. **(a)** The order-blindness argument we withdraw: for a *memoryless* unit, “A then B” and “B then A” produce identical sums (rust curve, second case overlaid dotted), so a threshold on that sum responds identically. This panel does not describe the present model, whose conduction delays and synaptic time constants make the summation history-dependent; a correlator built from the same units reaches $DSI = -0.79$ (Sec. 2.6). **(b)** Measured firing rates of T4a’s two correlator arms: the fast excitatory arm (Mi1) is suppressed by glutamatergic L1 while the slow inhibitory arm (Mi9) runs freely, a $37\times$ conductance imbalance. **(c)** Resulting direction selectivity with graded optic-lobe units. Magnitudes fall some $40\times$ short of the value at which experimental studies merely *select* a terminal as direction-tuned, and across 45 unsaturated operating points the two mirror pairs oppose at only 3, which is below chance: the sign is fluctuating, not small-but-real.

Both ingredients are present in the graph, and we verified them. T4a receives fast centred excitation (Mi1 +30.6, Tm3 +7.0) and slow flanking inhibition on opposite sides (Mi9 −7.4 at (−0.76, +0.26); Mi4 −3.4 at (+0.39, −0.60)), with four subtypes on four mirrored axes. The stimulus reaches them: L1’s response phase advances at −20.0 deg/deg against an expected 20.0, and this retinotopic travelling wave survives to T4a.

We implemented and measured six candidate remedies. Five produced nothing: per-type conduction delays (1.8–240 ms), per-neuron membrane time constants (5–200 ms), conductance-based shunting synapses, a forced high-conductance state ($g_{\text{tot}} = 105$, two orders of magnitude past the linear regime), and population pooling (per-cell DSI against per-cell input offset, $r = +0.043$ over 874 cells). The sixth, modelling the optic lobe as *graded, non-spiking* units, which is what photoreceptors and lamina monopolar cells actually are, was the only one to produce a measurable effect at all (Table 4), and it is small enough that its sign has to be checked rather than assumed.

Checking it requires sweeping the operating point rather than picking one, because the dominant influence on the number is not the wiring: graded units are capped at 200 Hz, and a unit against its ceiling reports $DSI \approx 0$ whatever its inputs do. We therefore swept bias, spatial period and temporal frequency (72 points), kept the 45 below 75% saturation, and report the largest effect among them. Two things follow. The magnitude is $|DSI| \leq 0.0121$, some $40\times$ below the value at which experimental studies merely *select* a terminal as direction-tuned. And the qualitative signature is absent: both mirror pairs oppose at only 3 of the 45 unsaturated points, where two independent coin flips would give about 11. A correlator’s sign does not depend on its operating point, so what we are measuring is a quantity fluctuating about zero rather than a small but real selectivity. Individual configurations can show apparent mirror opposition, but it does not survive the sweep.

One imbalance is measurable and specific, though the tests below show it is not sufficient to explain the failure: the fast arm is starved. Weighting T4a’s inputs by measured firing rate rather than synapse count gives excitation 189 against inhibition 558, a $3.0\times$ imbalance in drive; the realised synaptic conductances are $g_e = 0.0018$ against $g_i = 0.0667$, a ratio of $37\times$. The two quantities measure different things, rate-weighted input against the conductance it actually produces, and it is the second that sets

Table 4. Direction selectivity, at the strongest operating point that is not saturated. Graded units are capped at 200 Hz and a unit against its ceiling reports $DSI \approx 0$ whatever its inputs do, so we sweep bias, spatial period and temporal frequency (72 points), keep the 45 below 75% saturation, and report the largest effect among them — bias 1.0 mV, period 30° , 1 Hz. The magnitude is $41\times$ below the value at which experimental studies merely *select* a terminal as direction-tuned. The sign is not reliable either: both mirror pairs oppose at only 3 of 45 unsaturated points (7%), by pair T4a/T4b 5/45, T5a/T5b 15/45. A correlator’s sign should not depend on the operating point, so this is a quantity fluctuating about zero rather than a small but real selectivity.

Pathway	DSI		Mirror pair
ON	T4a −0.00431	T4b −0.00570	same sign
OFF	T5a −0.01208	T5b −0.01147	same sign

how far the excitatory arm can move the membrane. Mi1 sits at 2–9 Hz because L1 suppresses it (−77.3, glutamatergic and therefore inhibitory), while Mi9 runs at 78–130 Hz. The ON pathway is disinhibitory and, with uniform per-type gains, never rebalances. The OFF pathway, which does not depend on that disinhibition, conducts normally (T5a 58 Hz against T4a 2 Hz), an internal control indicating the deficit is pathway-specific rather than a global failure of the optic lobe.

Correcting that imbalance does not rescue selectivity. We rebalanced the arms two ways: presynaptic short-term depression (Tsodyks–Markram, $U = 0.3$, $\tau_{\text{rec}} = 0.5$ s), which depresses the tonically active inhibitory arm more than the sparsely active excitatory one; and transplanting the per-type gains of the optimized model of Lappalainen et al. [2024] onto the matching connections (571 of its 604 type-pairs align with our graph). Both *lowered* $|DSI|$. The $37\times$ ratio is therefore a real property of the wiring under a single global gain, but it is not the operative cause, and the hypothesis drawn from it, that optimization supplies relative gain along the ON pathway’s disinhibitory chain, is not supported by the transplant. We note the transplant is approximate: that model was fitted on a different reconstruction (fib25–fib19), so an unmatched gain is not evidence against the hypothesis, only against the simplest version of it.

Two further explanations were tested after the analysis above and are reported here because they are the two that a reader would reach for next.

The competing-input hypothesis. If the correlator is present but swamped, removing everything else should reveal it. We tested this in the RUNNING network rather than by isolated replay, which the caution at the end of this section explains cannot be trusted. Silencing every input onto T4a/T4b except Mi1 and Mi9, 14,594 of 24,944 edges, does not raise selectivity: over the 45 operating points unsaturated in both conditions, max $|DSI|$ moves from 0.0057 to 0.0020. Adding back Tm3, Mi4 or CT1 individually gives 0.0035, 0.0021 and 0.0034. T5a/T5b are never ablated and serve as a within-run control: their mean $|DSI|$ is 0.00442–0.00444 in all five conditions, so nothing leaked. The manipulation is electrically real: T4a’s rate ranges 63 to 97 Hz across conditions. The correlator fails even when it is the only thing connected.

Dendritic compartmentalization. A point neuron makes every synapse electrically identical, so a shunt divides the whole cell, whereas null-direction suppression is local to a branch. We split all 11,822 T4/T5 into a spiking soma and a passive dendrite coupled by an axial conductance, assigning each input by retinotopy alone: same column to the soma, different column to the dendrite. The rule reproduces the anatomy without naming a cell type (Mi9 97% distal, Mi4 98%) and is electrically effective (rates fall 193 $\rightarrow$ 151 Hz; unsaturated points rise 45 $\rightarrow$ 63). It does not rescue selectivity: max $|DSI|$ 0.0121 $\rightarrow$ 0.0128, mirror-pair sign below chance throughout. Two compartments cannot represent the ordering of arms along a real T4 dendrite, which is the obvious next refinement.

A related measurement bounds where the failure cannot be. At the tonic bias of 7.5 mV that makes an inhibition-dominated circuit conduct, the arms carry almost no time-varying signal at all: Tm3 and Mi4 sit at their ceiling 94.7% and 75.6% of the time with modulation amplitude 0.000 Hz, and Mi9 delivers 0.93 Hz out of a 200 Hz range. A correlator has nothing to correlate there. The operating-point sweep already excludes that regime, its filter removes every point above 3 mV, but the filter tests T4’s own saturation, not its inputs’, and those are different questions. Within the surviving range the two move together: as bias falls from 3.0 to 1.0 mV, arm modulation rises and max $|DSI|$ rises with it,

0.0073 $\rightarrow$ 0.0121.

That trend does not extrapolate, and the way it fails is worth recording because it is a trap in the measure itself. Extending the sweep below the range previously tested, down to zero tonic drive, max |DSI| keeps rising, to 0.0164. It is not signal. T4a fires 2.0 Hz there, and the quantity a correlator actually delivers, the rate DIFFERENCE between preferred and null, falls over the same interval, from 0.31 Hz at 0.75 mV to 0.060 Hz at zero. The ratio grows because its denominator collapses faster than its numerator. The mirror-pair sign agrees: across these points the pairs share a sign on 78% of them, where a correlator requires opposition and chance would give 50%. DSI is therefore interpretable only in a middle band: bounded above by saturation, which the sweep already controls, and below by the vanishing denominator, which it did not. We report the bounded figure, $|DSI| \leq 0.0121$, rather than the larger and emptier one.

What remains is a negative result without a mechanism, and we report it as such. Nineteen candidate explanations were measured and eliminated: pooling artefact (per-cell $r = +0.043$ over 874 cells); arm imbalance, by depression and by gain transplant; input dilution (the correlator supplies 80–86% of T4a’s input); unmodulated arms (Mi1 modulates at 125% and Mi9 at 89% of their means); CT1 saturation (12,601 edges ablated, no change); absent phase offset (measured mean 154° between arms); smeared spatial sampling (inter-column coherence 1.000); same-column sampling (Mi9/Mi4 sample one column away, 1.5% overlap); opposite-flank cancellation (flanks at 139° ; ablating one does not help); graded versus spiking units; tonic bias 0–7.5 mV; global optic-lobe gain $\times 1$ –32; gain targeted onto T4/T5 $\times 2$ –20; operating point, where the isolated correlator’s monotonic DSI-versus-rate curve does not reproduce in the network; the competing-input hypothesis, by in-network ablation; and dendritic compartmentalization. The geometry, the delays, the phase offset and the modulation are all present and measured; per-cell selectivity is not, though a population signal is (Sec. 2.10). Each of these interventions addresses a component that is not defective, which is why none of them helps. The operative cause lies in how the components are combined, and is identified in Sec. 2.7.

Two properties of the model, rather than of the connectome, account for a large part of the deficit, and both were found by building a correlator that works and then adding T4’s properties to it one at a time.

The cell is below threshold on its own inputs. Summed over the graph, T4a receives 46.8 units of excitation and 41.5 of inhibition per cell. Calibrating the same units in isolation, a target driven at the graded ceiling emits nothing at a weight of 30 and 24 Hz at 60: the threshold for any output at all falls at roughly 40–50. Supplying the hand-built correlator with T4a’s measured weights therefore leaves it silent, $R_{\text{pref}} = R_{\text{null}} = 0$. Under a single global W_{syn} calibrated on the SEZ taste circuit, the cell can only be made to fire by the tonic bias, which carries no stimulus information, so the selectivity we have been measuring is a ripple on an uninformative pedestal. T5a, at 58.9 against 31.3, clears the same bar, and reads $3.7\times$ the selectivity.

The graded transfer function removes the required nonlinearity. A correlator converts a phase difference between its arms into a rate difference, which requires an expansive nonlinearity. A graded unit emits clamp($(v - V_{\text{rest}})/\Delta, 0, 1$), a linear ramp between its clamps. Holding the recorded Mi1/Mi9 traces, the weights and the delays fixed and changing only the target’s transfer function gives DSI = -0.1781 spiking against -0.0168 graded, a $10.6\times$ loss; the latter is the value the intact network reports. Modelling the optic lobe as graded is correct for photoreceptors and lamina monopolars and was adopted for that reason, but applied to T4 it removes the operation the cell exists to perform.

Correcting both, spiking T4/T5, and optic-lobe gain in place of the tonic bias, raises the fraction of cells above the experimental selection threshold from 2.8% to 12.3%. It does not produce direction selectivity. The preferred direction remains set by the stimulus rather than by the wiring: across six spatial period and temporal frequency combinations, the fraction of T5a cells preferring one direction swings from 34% to 70%, crossing and reversing, and the residual response shows no delay tuning at all: varying the slow arm from 20 to 160 ms moves the directional signal only between 1.04 and 1.29 Hz. A correlator’s preference is a property of its geometry and cannot depend on the stimulus that probes it. We therefore do not describe the residual as weak direction selectivity; it is a stimulus-timing artefact, which also explains the below-chance mirror opposition reported above.

One methodological result follows from this work: isolated replay is not a valid control for this

question. Driving T4a cells with their own recorded inputs in isolation yields mean $|DSI| = 0.37$, which invites a causal account. The same procedure applied to a *single* input, where the geometry makes direction selectivity impossible, returns 0.12 rather than the required 0. The cause is that Mi1 is itself weakly selective ($|DSI| = 0.040$) and a threshold unit amplifies that residue. Isolated replay is not a clean control for this question.

2.7 Cancellation of the shunting nonlinearity

Direction selectivity requires a multiplicative interaction between the two correlator arms. In a conductance-based neuron the only available source of that interaction is inhibitory division: the membrane relaxes toward $v_\infty = (V_{\text{rest}} + g_e E_e + g_i E_i) / (1 + g_e + g_i)$, so inhibition divides rather than subtracts.

Measurement at the operating point shows the division is cancelled by its own consequence. Inhibition raises g_{tot} from 2.21 to 3.92, a factor of 1.77, and simultaneously lowers v_∞ from -23.5 to -43.4 mV, which raises the excitatory driving force ($E_e - v$) by a factor of 1.85. The net interaction is $1.85/1.77 = 1.056$. Across T4a, T4b, T5a and T5b the interaction does not exceed 11%, and its sign is inverted: inhibition slightly increases excitatory gain. Fitting the realised conductance trajectory confirms this independently. A purely additive model $v = a + b g_e + c g_i$ accounts for $R^2 = 0.865$ of T4a's membrane voltage, and adding the bilinear term $g_e g_i$ increases this by 2.2 percentage points.

The cell therefore sums its inputs. Summation is symmetric in its arguments and cannot recover temporal order at any wiring, which accounts for the failure of the nineteen candidate remedies listed above: each addressed a component that was not defective.

The interaction is governed by g_i/g_e rather than by g_{tot} . When the two conductances are comparable, the driving-force term cancels the shunt; the division survives only when inhibition substantially exceeds excitation.

A direct test separates the two components of inhibition and shows that the residual selectivity does not come from division at all. At $E_i = -70$ mV inhibition both hyperpolarises and shunts. Setting $E_i = V_{\text{rest}}$ removes the hyperpolarising component and leaves inhibition purely divisive, which is the condition under which the multiplicative account predicts the largest effect. Selectivity instead vanishes, and it falls off monotonically as E_i approaches V_{rest} from below:

E_i (mV)	-52	-60	-65	-70	-75	-80	-90
projected DSI	+0.0023	-0.0176	-0.0573	-0.0716	-0.0792	-0.0578	-0.0253
cells	3926	3325	2369	1872	1410	1017	608

The selectivity that survives is therefore produced by hyperpolarisation acting against the spike threshold, not by shunting division. A threshold supplies a coarse conjunction, since two inputs arriving together can cross it where neither does alone, and that is a poor substitute for a product, which is consistent with the magnitude observed. Two features of the sweep bound the claim. The included population shrinks as E_i falls, from 3926 cells to 608, so part of the variation is selection among cells driven above the inclusion rate rather than a change in selectivity itself. And the curve peaks at -75 mV against a published value of -70 mV, a difference within the standard errors (0.0064 and 0.0055), so the measured reversal potential is at or near the optimum and is not the parameter holding the model back.

2.8 Relaxing the point-neuron assumption

Every result above treats a neuron as one electrical node. That assumption is what makes the cancellation of Sec. 2.7 inevitable: if all inputs land on the same node, *where* an input arrives cannot be represented, and direction is exactly a statement about where. We therefore relaxed it. Each T4/T5 cell is split into a chain of three compartments coupled by an axial conductance, and every input is placed along that chain by its own retinotopic offset projected onto the cell's own correlator axis. No cell type is named in the placement rule; the axis is measured from the wiring, as in Sec. 2.6. Two controls accompany it: a *reversed* cable, distal and proximal exchanged, and a *shuffled* cable in which each cell keeps its inputs and the same number in each compartment, and only which input lands where is permuted.

parameterisation	values from	proj. DSI	behav.	gustatory	olfactory
(i) single global gain	fitted on taste	−0.0107	6/18	pass	pass
(ii) per-receptor, global	published physiology	−0.0356	4/18	pass*	fail
(iii) uniform $\times 2$, global	search	−0.0533	5/18	fail	n/a
(iv) $\times 2$ on vision only	search	−0.0716	6/18	pass	pass
(v) per-receptor, vision only	published physiology	−0.0500	n/r	pass	pass

Table 5. Parameterisations of synaptic conductance. Projected DSI is the population mean after projecting each cell’s response onto its own correlator axis; the raw population mean cancels across cells whose axes oppose. Behavioural column gives the number of eighteen metrics on which the model separates from its own command-matched control over thirty held-out environment seeds (40–69), all four rows run through the identical protocol on the same seeds. The column does not order the rows: selectivity rises monotonically down it while the behavioural count does not, and row (iv) does not exceed row (i). Row (iii) fails because bitter stimulation drives the proboscis motor pool to 168 Hz instead of suppressing it. Row (ii) passes the gustatory checks but degraded (*sugar 63.79 ± 9.31 against 73.58 ± 1.44 Hz, suppression 95% against 99%) and fails olfaction. Both global assignments break a circuit the frozen model preserves; neither regional assignment does. Rows (ii) and (v) apply the same published ratios and differ only in scope, which isolates the regional requirement from the value of the ratios themselves. Row (v) carries no fitted quantity at all and is marked n/r because its behavioural battery was not run: rows (i) through (iv) are mutually indistinguishable on that column, so a fifth measurement would not separate it.

Compartments buy an order of magnitude. Under a drifting grating, the separation between the T5 mirror pair (T5a minus T5b, measured per eye because the anatomical map mirrors them) rises from +0.005 as a point neuron to +0.064 with an ordered cable at an axial conductance of 8, with the vertical subtypes staying near zero on a horizontal grating. Individually T5a reaches +0.034 and T5b −0.035, against +0.001 and −0.004 as point neurons. This is the largest single improvement in per-cell selectivity we have found, and it is still some $15\times$ below the 0.5 at which a terminal would merely qualify as direction-tuned.

What it is not is a result about dendritic *order*. Four independent shuffled cables reach 0.047–0.050, reproducible to 0.003, against the ordered 0.064; and the reversed cable, at 0.086, is the best of the three. Most of what the compartments buy is therefore electrotonic: inhibition sitting electrically near the spiking node, which is the “on the path to the output” version of the shunt a point neuron cancels. Retinotopic sequence adds about 30% on top of that.

That statement is gain-dependent, and we report the dependence rather than the gain that flatters it. Raising the excitatory conductance globally by 1.25 – a single scalar applied to no particular cell type – inverts the ranking: the ordered cable gives +0.045, the reversed +0.039 and the four shuffles 0.011–0.013, which is the order a correlator predicts with sequence carrying most of the effect. The same scalar moves the point neuron’s separation from +0.005 to −0.029, and a gain of 1.5 to −0.043: an order of magnitude, *with a reversal of sign*. We draw the conservative conclusion. The mirror-pair structure is robust; the absolute direction the model prefers is not, and should not be read as a prediction.

Two further mechanisms were tested here and neither helped. (i) A *supralinear* dendrite, in which excitation is multiplied by a sigmoid of the compartment’s own voltage, as an NMDA-type Mg block would be — the preferred-direction enhancement that accompanies null-direction suppression in *Drosophila* [Haag et al., 2016, Gruntman et al., 2018]. Against a control with the same mean amplification and no voltage dependence, the difference is −0.0425 against −0.0428: the nonlinearity contributes nothing that the extra drive does not. (ii) Replacing the slow arm’s fixed conduction delay with a one-pole low-pass filter, which is the shape a biological slow arm has. Selectivity falls at every time constant tried, peaking at 0.039 where the filter’s time constant equals the delay it replaces, against 0.064 for the delay. The sharp temporal edge is what the correlator multiplies against, and smearing it costs about 40%.

2.9 Comparison of conductance parameterisations

The preceding result makes the assignment of synaptic conductances the study’s operative variable. We compare three ways of setting it, together with the uniform scaling that motivates the third.

(i) *A single global gain*. This is the frozen model. Direction selectivity is $40\times$ below the experimental

selection threshold and the model does not separate from the control on any behavioural measure.

(ii) *Per-receptor conductances from published physiology.* Neither value was chosen with reference to any result reported here. GABAergic conductance is $2.5\times$ cholinergic in two unrelated preparations: Su and O'Dowd [2003] record quantal events for both transmitters in the same cells, and Groschner et al. [2022] fit adult T4 in vivo. We take $g_{\text{GluCl}}/g_{\text{ACh}} = 0.92$ from the latter and note that it is a fit rather than a measurement, no direct *Drosophila* GluCl conductance having been published. Applying both ratios raises g_i/g_e from 0.713 to 1.323 and improves selectivity $3.3\times$, with the absolute rate difference rising from 1.09 to 2.97 Hz. It nonetheless fails the olfactory constraint. DNp01 falls to 0.00 Hz in both the odour and control arms: raising inhibition brain-wide silences the neuron outright, so no baseline discharge remains for the odour to suppress and the measured difference is zero by construction rather than by loss of the pathway. LHN falls from $+14.38$ to $+4.16$ Hz and DNp02 from $+12.12$ to $+0.05$ Hz. This is the informative result in the comparison. The values are not ours and were not selected to produce any outcome reported here, yet applying them uniformly breaks a circuit that the cruder single-gain model preserves. Correct per-receptor ratios are therefore necessary but not sufficient: the assignment must also be regional, which is a property of this brain that receptor measurements cannot themselves supply.

(iii) *Uniform scaling, applied globally.* Doubling every inhibitory synapse improves selectivity but inverts the gustatory response: bitter stimulation drives the proboscis motor pool at 168 Hz instead of suppressing it, and the dose-response relation degrades to noise. The inversion occurs because 557,080 of 1,099,675 inhibitory edges (51%) terminate on inhibitory neurons, so uniform scaling amplifies disinhibition equally and the sign of the circuit reverses.

(iv) *The same factor restricted to the visual populations.* The gustatory result is numerically unchanged to the reported precision (sugar 73.58 ± 1.44 Hz, bitter 0.00, 99% suppression, six of six acceptance checks), because MN9 receives 0.00% of its input synapses from the scaled populations and a regional parameter has no path into the circuit. Olfaction is preserved, including the readout that is structurally exposed: DNp01 draws 20.5% of its input from the scaled populations and holds at -17.02 ± 7.15 Hz against a baseline of -18.13 ± 2.48 Hz. The one cost is DNp02, which falls from 12.12 to 2.65 Hz.

Rows (iii) and (iv) differ only in the region to which an identical factor is applied. The comparison establishes that a single global synaptic gain is not merely an approximation but an inconsistency: the optic lobe requires a conductance ratio the subesophageal zone does not tolerate. Groschner et al. [2022] report inhibitory-to-excitatory conductance ratios of 7.7 and 10.9 in adult lobula plate neurons, so the factor of two used here is conservative relative to the published requirement.

(v) *Published per-receptor conductances restricted to the visual populations.* Rows (ii) and (iv) vary two things at once: (ii) supplies measured ratios but applies them brain-wide, and (iv) applies a factor obtained by search but restricts it to vision. The combination the argument implies, measured ratios restricted to vision, completes the design. It uses the ratios of row (ii) and the scope of row (iv), and introduces no quantity of our own. Projected selectivity reaches -0.0500 against -0.0107 frozen, the gustatory result is numerically unchanged (sugar 73.58 Hz, 99% suppression, six of six checks) and olfaction is preserved (lateral horn $+12.42$ Hz, DNp01 -17.55 Hz), where the identical ratios applied globally silence DNp01 altogether. Scope alone therefore accounts for the difference between failure and preservation, with the values held fixed. Row (iv) reaches higher selectivity, but by a factor chosen to do so; row (v) reaches most of it with published numbers throughout, and is the configuration we would defend as the model rather than as a tuning.

The accessible range is bounded. Selectivity was measured across the parameters that do not entail per-cell-type fitting: the inhibitory reversal potential from -52 to -100 mV, the regional inhibitory scale from 1 to 3.5, per-receptor ratios global and regional, and the optic gain from 4 to 64, this last raising selectivity monotonically but shallowly (-0.0245 to -0.0596) at firing rates well above physiological. The largest value found anywhere was -0.0792 , at $E_i = -75$ mV with the regional scale at 2. Every configuration in this space therefore remains close to an order of magnitude below the threshold at which experimental studies classify a terminal as direction-tuned, which makes the shortfall a measured property of the accessible space rather than a consequence of the particular values we report.

2.10 Where the directional signal is read

Direction selectivity above is reported per cell, which is what the experimental literature reports for single terminals. Behaviour does not run on a single terminal. It runs on whatever the steering pathway reads, and the choice of recording site turns out to matter more than any parameter examined so far.

The steering readout used throughout is DNa02, which the connectome contains as one cell per hemisphere. That pair carries a standing left–right difference of 20 to 45 Hz, an asymmetry of this particular brain rather than a response, on which the directional modulation is 0.5 to 2.5 Hz and does not hold its sign across temporal frequency. The instrument cannot resolve the quantity.

The horizontal system is the optomotor output stage in the animal: HSE, HSN, HSS and H2, four cells per hemisphere whose function is to pool thousands of T4/T5 terminals across the eye, and where the optomotor literature records. Reading there is an electrode placement and changes nothing in the model. To separate a genuine motion response from an asymmetry of the sampling lattice, every input onto T4/T5 is zeroed in a control condition, which removes the motion detectors while leaving the stimulus and the rest of the brain untouched. Mirroring the retina is not a control here: for a drifting grating it is algebraically the reversal of drift that the measurement already compares.

drift (Hz)	0.5	1	2	3	4	6	8
intact	29.5	19.6	25.0	17.3	16.4	13.4	11.8
T4/T5 ablated	3.4	16.1	0.9	0.4	5.0	7.0	0.6
difference	26.1	3.5	24.1	16.9	11.4	6.4	11.1

At five of seven frequencies the control lies between 0.4 and 3.4 Hz against an intact response of 12 to 30 Hz, so the frozen model carries a directional signal at the optomotor output stage that is mediated by T4/T5 and is an order of magnitude larger than anything DNa02 resolves. Two qualifications belong with it. The control is not a constant that can be subtracted, reaching 16.1 Hz at 1 Hz drift and 7.0 at 6 Hz. And the directional term rides on a standing offset of 70 to 100 Hz rather than crossing zero, where an animal's would swing symmetrically about it.

The selected parameter is the worst at this site. Measured as intact minus ablated at the horizontal system, the frozen model gives +26.1, +24.1 and +16.9 Hz at 0.5, 2 and 3 Hz drift. The regional scale of 2, which raises per-cell selectivity sevenfold and which the 137-point search selected, gives −9.5, −1.5 and −0.25 at the same frequencies. The two quantities are not merely different; across the scale from 1 to 3 they move against each other. The search maximised per-cell selectivity because that is the published quantity, and in doing so removed the population signal that a steering pathway would read.

Selectivity and behaviour were measured on different models. A graded unit in this simulator is a clamped linear ramp and has no threshold. The reversal-potential sweep of Sec. 2.7 showed the residual selectivity here to be threshold-derived, so a model without thresholds should have none, and it does not: the T4/T5-dependent signal at the horizontal system is +24.1 Hz with spiking T4/T5 and −0.12 Hz with graded. The closed-loop agent runs graded by default and every behavioural run used it, while every selectivity measurement, including the search, exempted T4/T5 from the graded mask. At the same regional scale projected DSI is −0.0716 spiking against −0.0336 graded, so the reported figure describes a configuration that never ran in the loop. The behavioural battery below therefore repeats every condition in both transfer functions.

2.11 Where the signal is read on the motor side

Sec. 2.10 located a directional signal at the horizontal system that the behavioural measures do not reflect. The gap is in the readout rather than in the signal.

The steering readout is almost blind. DNa02 draws 2.3% of its input weight from visual populations and 0.34% from the horizontal system, the remainder arriving from the central brain. It is excitation-dominated ($g_i/g_e = 0.56$) and fires at 140 and 400 Hz on the two sides against a refractory ceiling of 454 Hz, so it is both weakly connected to vision and close to saturation, where a neuron has little dynamic range in which to represent anything. DNp01, the neuron the olfactory result moves, is the opposite on both counts: 20.3% visual, inhibition-dominated at 1.81, with a baseline that an input can modulate.




Figure 4. The directional signal reaches a descending neuron, but not the one the decoder reads. **(a)** Share of each descending neuron’s input arriving from visual populations. DNa02, which the motor decoder uses for steering, receives 2.3%; several DNp types receive 50–80%. **(b)** Firing-rate census of all 1,305 descending neurons in the running loop: 79.8% never fire, and the 19 above 200 Hz include DNa02, against a refractory ceiling of 454 Hz. **(c)** The T4/T5-dependent left-right difference, measured as intact minus ablated, at three recording sites.

The classic pathway is not in this reconstruction. The optomotor pathway described in the literature runs from the horizontal system to DNa02. Here the horizontal system sends 14.5% of its output to descending neurons, of which 75 of 7,297 units of signed weight reach DNa02. Its strongest descending target is DNp15, at $4.7\times$ that weight.

And DNp15 carries the signal. Measured as intact minus T4/T5-ablated, the left-right difference at DNp15 is +19.0, −4.0, +32.0, +26.0, +11.5 and −20.5 Hz at 0.5 to 6 Hz drift: substantial and positive at four of six frequencies, larger where present than at the horizontal system itself, and with sign inversions at 1 and 6 Hz that the same ablation control shows to be artifact-dominated frequencies. The visual pathway therefore delivers a directional signal to a descending neuron. The decoder reads a different one.

The published characterisations of both neurons agree with these numbers, and they identify the mismatch as a category error rather than a tuning problem. DNa02 steers goal-directed walking: it is targeted directly by central-complex output and lies two synapses from the head-direction system, being recruited when a heading is compared against an internal goal [Rayshubskiy et al., 2025]. That is what a neuron drawing 97.7% of its input from the central brain should be. DNp15, also called DNHS1, is the optomotor neuron: it receives strong input from the horizontal system and projects to the neck motor neuropil, and it is one of the descending cells encoding self-motion from optic flow [Suver et al., 2016, Namiki et al., 2018]. The model has been delivering optic flow to the navigation pathway and asking whether vision steers.

Because this environment’s yaw control rotates the viewpoint, and DNp15 rotates gaze through the neck, the correspondence is if anything closer than the incumbent’s. We therefore tested the readout rather than assuming it, selecting it on published function and on connectivity rather than on which cell carried the signal we were looking for.

Changing only the readout. With spiking T4/T5, an optic gain of 16 and the same thirty held-out seeds, the two readouts differ in their coupling to the retina and not in task performance:

yaw from	separating	in the model’s favour	vision–steering r
DNa02 (navigation)	5/18	1/15	−0.043
DNp15 (optomotor)	7/18	4/15	−0.158

Paired by seed, the vision–steering correlation moves -0.12 ± 0.10 and mean yaw magnitude falls 0.43 ± 0.19 ; no task metric moves significantly, and 7/18 against 5/18 is not a separation. Mirroring the retina, which reverses every horizontal optic-flow signal while leaving rates, contrast statistics and wiring untouched, abolishes the difference: −0.158 returns to −0.034, a paired change of $+0.124 \pm 0.102$, against the DNa02 baseline of −0.043. The coupling is therefore retinotopy-dependent, which an artifact

of wiring or of motor statistics would not be.

It does not reverse. A clean optic-flow relationship should change sign under a mirrored retina rather than fall to zero, so what is established is dependence on vision and not the sign structure of an optomotor response. Two of the three mismatches, the transfer function and the optic gain, are null in isolation over the same seeds (6/18, 4/18 and 5/18); the third moves the sensorimotor coupling and not the behaviour.

The output stage as a whole. Of 1,305 descending neurons, 1,041 (79.8%) never fire during closed-loop operation, 134 sit in the physiological 1–50 Hz band, and 19 exceed 200 Hz. The median is zero. A model whose motor output is mostly silent and whose active minority runs near saturation constrains any behavioural claim independently of what the sensory side computes.

2.12 A third sensory channel, and why it also fails

The model had vision, smell and taste. Walking into a wall produced no afferent activity anywhere, which makes collisions, the one behavioural measure that moves under every manipulation reported here, a consequence the agent cannot sense. The connectome carries 2,674 mechanosensory neurons that go unused. The environment supplies contact on the same terms it supplies odour sources, so the omission is ours rather than the engine's.

The obvious target is the wrong one, and the reason is instructive. FAFB contains 1,113 eye bristles and 304 head bristles lying one synapse from 109 descending neurons, which are DNg15, DNg35, DNg48, DNg84 and DNg85. The DNg group is largely grooming, and that is correct biology: deflecting a fly's head bristles elicits head grooming rather than avoidance. Driving wall contact into bristles would ask the model to groom on impact. Flies wall-follow by antennal contact, and the 486 wind and gravity afferents lie two synapses from DNa02 against roughly six for the visual pathway, on a route that does not depend on the multiplicative interaction this neuron model fails to produce.

Contact is derived rather than read: an agent commanding forward motion it does not achieve is against something, and the side follows from the slide, since a body pushed into an angled surface slides along it. Contact reaches DNa02 and shifts it by roughly 3 Hz, but the shift does not reverse with the contacted side (−3.12 against −3.66 Hz, sign consistent on two of three seeds) on a standing asymmetry of −115 Hz. Summed two-hop signed weight explains why: the left antenna reaches DNa02_L 1.35× more strongly than DNa02_R, and the right antenna also favours DNa02_L, by 1.26×. Both antennae drive the same side, so the pathway is not differentially lateralised at two synapses and cannot produce a side-dependent turn however the afferents are driven. Whether the real reflex runs through the ventral nerve cord, which this reconstruction does not contain, we cannot say.

Three sensory channels now give three distinct failures: vision computes a signal that reaches a descending neuron the decoder does not read, touch reaches the decoded neuron but without lateralisation, and smell reaches it lateralised and with the predicted sign, which is the one that works.

2.13 Closed-loop behaviour

Every condition was run through the identical protocol on the same thirty environment seeds (40–69), none used while selecting parameters, each against its own command-matched AR(1) control and paired by seed, because seed-to-seed variance here is far larger than the differences under test. A metric counts as separating when the paired difference has a 95% confidence interval excluding zero. The second column counts only metrics with a defined better direction: three of the eighteen (path length, and mean yaw and forward magnitude) describe what the agent did rather than how well, and are excluded rather than assigned a sign.

parameterisation	separating	in the model's favour	vision-steering r
(i) frozen	6/18	2/15	−0.090
(ii) per-receptor, global	4/18	2/15	+0.017
(iii) uniform, global	5/18	3/15	−0.031
(iv) regional	6/18	3/15	+0.161

No parameterisation separates from its control by more than the frozen model does. Restoring the multiplicative interaction therefore raises per-cell direction selectivity sevenfold without producing

task-level behaviour.

The signal does not reach the readout. Because the directional signal at the horizontal system varies by an order of magnitude and by sign across these configurations (Sec. 2.10), the transfer function and the regional scale were crossed and all four cells run on the same thirty seeds.

T4/T5	scale	HS signal (Hz)	separating	vision-steering r
graded	1	-0.12	6/18	-0.090
graded	2	-	6/18	+0.161
spiking	1	+24.1	4/18	-0.066
spiking	2	-1.5	6/18	+0.144

The two spiking conditions differ by an order of magnitude and in sign at the optomotor output stage and return 4/18 and 6/18. Compared directly against the graded model on identical seeds, the spiking model covers more ground ($+4.63 \pm 3.07$ tiles), collides less (-3.48 ± 3.08 per 1,000 tics) and displaces further ($+108 \pm 106$), but the vision-steering correlation moves $+0.02 \pm 0.10$. The regional parameterisation produced the same signature against the frozen model. Two mechanically unrelated interventions leaving one fingerprint, in metrics of locomotion and not of vision, is more consistent with altered motor statistics than with sensory control.

A retraction that follows from the grid. The vision-steering correlation tracks the inhibitory scale in both transfer functions, near -0.07 at 1 and $+0.15$ at 2, and does not track the directional signal, which the scale of 2 abolishes. It is therefore not a clean index of optic flow reaching the steering readout, and we withdraw it as behavioural evidence of vision on its own. It is retained in the tables as a measured quantity.

Sec. 2.11 qualifies rather than reverses this. Holding the inhibitory scale fixed and varying only the readout, the same measure is abolished by mirroring the retina, which makes it retinotopy-dependent in that comparison. Both results stand: the measure is sensitive to the inhibitory scale and to retinal organisation, and reporting it as evidence of vision requires the mirrored control beside it rather than the number alone.

This is the outcome the power analysis in Sec. 4.6 predicts. Health collected carries a standard deviation of roughly 13 units across seeds, so thirty seeds resolve differences of about 9 units, and two executions of one identical configuration have returned means of 3.2 and 11.2. An open-loop measurement that moves sevenfold is simply not visible through this much environmental variance at any seed count we can afford: detecting a 5-unit difference at 80% power would require $n \approx 106$. We report the negative result rather than a subset of metrics chosen after the fact, and note that it constrains what any closed-loop connectome study can claim from episode-level measures.

A vision-dependent advantage over matched chance. The comparisons above change the stimulus between arms, and in closed loop that confounds every measurement: an arm that behaves differently goes somewhere different and sees something different. The comparison between the model and its own command-matched control does not have that defect, because both arms receive identical conditions and differ only in the source of the command. In the configuration that applies every correction reported here (spiking T4/T5, optic gain 16, yaw decoded from DNp15, antennal contact), that comparison separates, and it separates only while the model can see.

model minus matched control	intact retina	mirrored retina	frozen retina
health collected	$+12.67 \pm 10.80$	-0.53 ± 7.37	-5.60 ± 11.57
tics survived	$+37.90 \pm 36.60$	-5.47 ± 31.88	-33.30 ± 48.33
tiles visited	$+3.93 \pm 3.05$	-0.57 ± 3.71	$+0.47 \pm 4.22$
stuck fraction	-0.07 ± 0.06	$+0.01 \pm 0.08$	-0.06 ± 0.08

With an intact retina the model beats the matched control on four measures, over thirty held-out seeds. Mirroring the retina, which reverses every horizontal optic-flow signal while leaving firing rates, contrast statistics and wiring untouched, abolishes all four. Freezing it on a single frame, which removes optic flow while preserving the mean image, abolishes them too. The two controls break vision

in different ways and agree, so the conclusion does not rest on either manipulation alone. Absolute health collected is graded across the conditions, 24.80 intact, 16.80 mirrored, 14.27 frozen.

The claim is narrower than it may appear, and we state its limits. The dissociation lies in the advantage over the matched control, not in absolute performance: the model’s own decline without vision is directionally consistent in every comparison and individually significant in none (frozen minus intact, health -10.53 ± 13.31). The control is matched per seed to the model’s own command statistics and itself improves under mirroring, from 12.13 to 17.33 health, which means that without usable vision the model’s commands become ones a random agent does better with. What is established is therefore that the model’s command structure is advantageous only while it can see. What is not established is that any recognisable fly reflex underlies it: the optomotor drum shows no directional response, looming shows none, the fixation cells are silent, and klinotaxis is absent.

The full visual field, in a second arena. Every result above samples at most 64% of the lattice, because one planar view cannot cover two eyes of 170° splayed $\pm 40^\circ$. We removed that limit by rendering three views (Methods), so that all 4,541 columns sample the scene, and at the same time replaced the arena’s dark ceiling with open sky. The brain and decoder are unchanged. Sixty held-out seeds (40–99) were run per condition, in two blocks; the second was added because the first was underpowered, is reported combined, and nothing was selected on either.

	model minus matched control	full eye	mirrored	frozen	original eye
health collected		$+8.67 \pm 6.99$	$+3.67 \pm 5.71$	$+0.60 \pm 6.42$	-0.87 ± 6.08
tics survived		$+31.77 \pm 29.16$	$+11.73 \pm 27.02$	-4.75 ± 26.14	$+0.18 \pm 22.33$
tiles visited		$+3.83 \pm 2.42$	$+1.28 \pm 2.27$	$+3.02 \pm 2.64$	$+2.22 \pm 2.53$
stuck fraction		-0.06 ± 0.06	-0.05 ± 0.05	-0.06 ± 0.05	-0.04 ± 0.06

With the full eye the advantage in health and survival returns, with the same cost in damage taken ($+7.87 \pm 6.73$), and mirroring or freezing the eye removes it. With the original 130° view in the same arena there is no such advantage, so in this arena the effect requires the wider field. Two measures do not share the dissociation: exploration and stuck-fraction survive a frozen eye, so they are not evidence of vision here. The limit stated above also holds: the model’s own health is ordered as expected (21.2 full eye, 19.3 mirrored, 19.3 frozen, 16.3 original eye), but no model-versus-model difference in health or survival is individually significant (full minus original eye, health $+4.93 \pm 6.23$).

With the interface corrected. Every behavioural result above was measured through the interface whose errors Sec. 2.17 and the Limitations describe, including an agent that walked backwards on 97% of tics. We re-ran the full-eye comparison with all of them corrected: the same brain configuration, the same sixty held-out seeds and controls, in the sky arena and in an arena rebuilt so its walls carry structure at the fly’s resolution (Methods). The fly-arena conditions were then repeated on sixty *further* seeds (100–159), and the table reports all 120; the replication is discussed below.

	model minus matched control	sky: intact ($n = 60$)	fly: intact ($n = 120$)	fly: mirrored ($n = 120$)	fly: frozen ($n = 120$)
health collected		-4.27 ± 6.79	$+5.67 \pm 4.88$	-0.40 ± 4.72	$+3.27 \pm 4.41$
tics survived		-8.18 ± 26.56	$+21.01 \pm 19.31$	-7.11 ± 18.32	$+13.65 \pm 17.93$
tiles visited		-0.15 ± 1.76	$+1.34 \pm 1.50$	$+1.09 \pm 1.46$	$+2.93 \pm 1.38$
stuck fraction		$+0.02 \pm 0.05$	$+0.02 \pm 0.04$	-0.09 ± 0.04	-0.09 ± 0.04

In the sky arena the advantage is gone: the intact model does no better than chance on any measure, and head to head it covers 3.52 ± 2.23 fewer tiles and is stuck 0.09 ± 0.06 more than the mirrored one. In the fly arena the health and survival advantage returns with the profile it had before, including more damage taken ($+8.63 \pm 6.17$), and a mirrored retina removes it; a frozen retina only weakens it, so the second control is ambiguous rather than clean. The earlier result was therefore partly an artifact of its errors: reversing with a reversed turn sign steers the direction of travel toward the more active side, and the agent faced away from the wall it was pressed against. With correct geometry the model’s vision instead steers it toward large dark structure — it is stuck more than the mirrored model in both arenas — which is the attraction a fly shows to dark vertical objects, and maladaptive in a maze. What

survives is a vision-dependent advantage over chance, confined to a world whose structure the animal can resolve.

How well it replicates. The fly-arena result was first measured on seeds 40–99 and gave $+7.53 \pm 6.66$ health and $+34.00 \pm 25.27$ tics. Sixty fresh seeds, run identically, give $+3.80 \pm 7.16$ and $+8.02 \pm 29.04$ — half the size, and neither interval excludes zero on its own. Pooled over all 120 the advantage stands ($+5.67 \pm 4.88$, $+21.01 \pm 19.31$, both excluding zero) with the controls behaving as before, but it is about a quarter smaller than the first block implied, and the first block was evidently the favourable half. We report the pooled estimate throughout and note that a single block of sixty seeds was not enough to pin this effect to better than a factor of two.

Applying every later improvement at once removes the vision dependence. Three changes developed after these tables — T4/T5 as three-compartment cables (Sec. 2.8), a neck that turns the eyes up to 20° on the same yaw command, and odour delivered from the whole level rather than gated on what is visible — were run together on the same sixty seeds, so the comparison is paired and isolates them.

The task score does not move: health -0.13 ± 6.85 against the plain configuration. The movement is worse on every measure of control we record: turn-command chatter $+0.149 \pm 0.015$, collisions $+5.11 \pm 2.70$ per 1000 tics, tics between collisions -18.2 ± 15.7 , and the correlation between retinal left–right luminance and the turn command -0.034 ± 0.027 . The only measure that improves, ground covered per unit walked ($+1.31 \pm 0.76$), is the one the mirrored and frozen controls have always scored well on.

The control is what matters. In the plain configuration a mirrored retina abolishes the advantage over chance (-0.67 ± 7.70 health, against $+7.53 \pm 6.66$ intact), which is the basis for calling that advantage vision-dependent. With all three applied it abolishes nothing: intact minus mirrored is -0.50 ± 7.73 , and a model whose retina is *frozen on a single frame* beats its matched control outright ($+6.53 \pm 5.97$ health, $+30.55 \pm 22.79$ tics). The agent scores slightly higher and no longer requires its eyes.

Running the odour flag alone on the same seeds separates the causes and rules out the one we expected. Whole-level odour by itself produces no advantage in any arm: intact falls to $+3.80 \pm 6.73$ and ceases to be significant, mirrored sits at $+0.40 \pm 6.51$, and the frozen arm — a blind model with odour reaching every medkit at five times the previous strength — scores -0.07 ± 6.56 . Strong odour does not find food here, and the reason is structural rather than surprising: the olfactory input is bilaterally symmetric by construction, a scalar concentration with every azimuth discarded, since a direction-carrying odour channel would be the confound the experiment exists to avoid. A scalar cannot steer, so strengthening it adds drive without information, which is also why the intact model does *worse*. Odour is therefore not sufficient. Neither is the neck, run alone on the same seeds: $+2.67 \pm 7.02$ intact, $+3.67 \pm 7.28$ mirrored, $+0.53 \pm 6.64$ frozen. Each change *on its own* costs the advantage — $+7.53$ falls to $+3.80$ with odour and to $+2.67$ with the neck — without producing the broken control; only together do they restore performance, by a route that survives both mirroring and freezing. That is an interaction rather than a sum, and the compartment cables are the one component not yet run alone. We report the dissociation as established and its cause as open. The behavioural tables above retain the vision-gated olfaction regardless, since the whole-level variant changes the task without improving it.

One methodological consequence generalises. With a movable head, freezing the frame is no longer a no-vision control: the frame is held, but the eyes continue to turn across it, so the lenses sample a moving window of a static scene and the model receives self-generated visual motion. Any frozen-frame control must also pin the head, or freeze the sampled per-lens luminance rather than the frame. The frozen arm quoted above carries this confound.

Doubling the seeds does not rescue the model-versus-model contrasts, which is the limit named above. Intact minus mirrored on health is $+3.50 \pm 4.63$ at $n = 120$, against $+3.00 \pm 6.18$ at $n = 60$: the interval narrowed as $\sqrt{2}$ predicts and still contains zero. Reaching significance at this effect size needs roughly 210 seeds, and we record the number rather than running toward it, since the dissociation the paper rests on is between each model and its own matched control, not between models.

Three controls exclude the possibility that the gain in selectivity follows from reduced drive alone. Matching the firing rate by lowering gain instead of raising inhibition yields -0.0152 against -0.0716 . Applying the identical factor to non-visual inhibitory synapses yields -0.0102 against an

unmanipulated -0.0107 . On the 1,869 cells driven above threshold in both conditions, projected selectivity moves from $+0.0001$ to -0.0720 , excluding selection among a shrinking population.

The chosen scaling is not a sharp optimum. We initially described it as narrow, on the grounds that selectivity collapses at a scaling of four. A structured search over 137 parameter combinations measured the local curvature directly and does not support that reading. Holding the score within 20% of its maximum admits scalings from 1.67 to 2.81, a band $\pm 25\%$ about the chosen value on a three-stimulus mean; the collapse at four is the far shoulder rather than the local shape. That search found no feasible configuration superior to (iv). A scaling of 2.5 raises the ratio further but not the absolute rate difference, which is maximal at 2.0 (3.97 Hz averaged over three stimuli) and reverses on both held-out stimuli. Scaling the complementary non-visual population has no effect at any value between 0.5 and 3, despite rescaling 467,082 edges, which is a further specificity control. Every higher-scoring point the search located lay outside the feasible region, failing the olfactory constraint in the manner described for row (ii).

2.14 Bar fixation, and a control that passes

The two visual computations tested above are both computations over temporal order, which is the thing this model loses. Fixation is not: which side an object lies on is a spatial fact, and the retinotopy is verified end to end (Sec. 2.17). It is also, with the optomotor response, among the most robust behaviours an insect has. We therefore built the arena that behaviour is measured in — a 72-sided cylinder of radius 512 map units, uniformly bright, carrying one dark bar 20° wide, with the fly at the centre — and a twin cylinder with no bar at all. At constant distance the engine’s light falloff is constant, so the bar is the only azimuthal feature in the world. The fly is tethered: its forward and lateral commands are recorded and discarded and only its yaw reaches the body, so the bar’s retinal position is a pure integral of what the brain commanded. We report $F = \langle \cos \theta \rangle$, the mean cosine of the bar’s azimuth, which is $+1$ for a bar held ahead and -1 for one held behind.

The model locks the bar at a stable azimuth and holds it. Over 20 held-out seeds and 1200 scored tics the mean resultant length is 0.965, so once settled the bar stays within a few degrees for the rest of the episode, while the fly turns at 1.3° per tic and is turning on 78% of them.

It is not fixation, and the controls rather than the effect are what show it. A model whose retina is frozen on a single frame locks *more* tightly (0.980) and scores higher, $F = +0.299 \pm 0.309$ against $+0.128 \pm 0.278$ for the intact model and -0.051 ± 0.333 mirrored. In the blank cylinder, which contains nothing to fixate, the score is $+0.265 \pm 0.307$ — indistinguishable from the arena with a bar in it. The azimuth the intact model holds is 83° off centre on average, not ahead.

The decoder accounts for all of it. The yaw channel subtracts a 3 s running baseline (Sec. 2.11), so a sustained command is removed and the loop’s only fixed point is zero net turning; the heading then freezes wherever it happens to be, and because the bar is fixed in the world a frozen heading is a fixed bar azimuth. A closed-loop equilibrium measure cannot distinguish an animal that fixates from one that has stopped.

The open-loop version has the power the equilibrium measure lacks. Turning the body at a constant 0.5° per tic walks the bar twice round the eye in one episode, and the commanded yaw binned by bar azimuth is the position response directly. Each arm is run at $+0.5$ and -0.5 , because a response to the rotation is common to both and reverses with it while a response to position does not. The restoring slope, in degrees of command per degree of bar offset across the frontal $\pm 90^\circ$, is $+0.0003$ and -0.0008 for the intact model: it reverses with the rotation, so it is a response to being turned. Its magnitude is smaller than the same measurement in the cylinder with *no bar* (0.0012) and a third of the value a *frozen retina* produces (0.0024). The bar-dependent component of the turning command is therefore below what the apparatus generates on its own, against a command whose standard deviation is 1.8° per tic; the binned curves scatter between -0.50 and $+0.59$ with no systematic shape.

One methodological point generalises beyond this experiment. The four seeds returned byte-identical commands, because the model is noiseless, the cylinder is fixed, and the rotation is imposed: nothing in an open-loop drum episode depends on the seed. The measurement is exact rather than noisy, the scatter in the binned curve is history-dependence in the decoder rather than sampling error, and a confidence interval computed across seeds would be fiction. In a deterministic arena the unit of

replication is the stimulus — rotation direction, rate, bar width — and not the random seed.

2.15 A heading signal, and the wind that would use it

Asked how a fly finds its way, the answer is mostly not by seeing food. It stabilises course on optic flow, avoids collisions by expansion, and holds a direction with a compass: the ellipsoid body carries a ring attractor whose EPG bump encodes heading, anchored to landmarks and updated by self-motion. Two of this project's failures need not apply — direction selectivity and looming are computations over temporal order, which this model loses, whereas a compass is a persistent state. Every component is in the connectome: 51 EPG, 42 PEN, 20 PEG, 278 ER, 42 Delta7.

It is not there. Fitting each cell's preferred heading on the first half of an episode and decoding the second gives a circular correlation of $+0.129 \pm 0.277$, *below* the $+0.211$ scored by the same decode against a deliberately mismatched heading trace, and no better than 51 neurons drawn at random ($+0.075$). That last control is the one the measurement needs: in a fixed arena heading determines what is on the retina, so any visually driven cell correlates with heading and decoding one demonstrates nothing.

Odour cannot substitute. The olfactory input here is bilaterally symmetric by construction — a scalar concentration with azimuth discarded, so that a direction-carrying chemical channel cannot become the confound the study exists to avoid — and a scalar cannot steer. Measured, whole-level odour at five times strength leaves a blind model at -0.07 ± 6.56 against its matched control.

A fly solves exactly this problem with a different sensor: in turbulent air an odour gradient points nowhere, so the animal uses odour as a gate and airflow as the steering cue, casting across the wind and surging upwind on contact with a plume. We added that as physics rather than as a hint: a source's odour reaches only a fly downwind of it, and the flow direction arrives as a left-right difference at the Johnston's organ afferents already present in the connectome and already driven by the touch channel. Clamping the antennae, flow from the right against flow from the left moves DNp15's left-right differential by about 8 Hz with the same sign in three seeds of three — a larger and more consistent response than either antennal contact or the optomotor drum produces.

It does not find food: against a matched random agent, whole-level odour gives $+3.80 \pm 6.73$ health and adding wind $+1.40 \pm 9.25$. But health is the wrong measurable for an animal with no circuitry for collecting medkits, and the movement tells a different story. Paired on the same seeds, with the same odour, a steady wind reduces the net rotational bias from 0.167 to 0.128 (-0.039 ± 0.030) and raises path straightness from 0.462 to 0.562 ($+0.100 \pm 0.074$): a quarter less circling and a fifth straighter walking, from a non-visual cue, with nothing in the model changed. A fly stabilises its course on optic flow; this model cannot, and airflow partly substitutes for the sense that fails.

It is stabilisation and not steering, which the control establishes. Measuring how tightly heading concentrates relative to the wind gives a mean resultant of 0.639 with airflow and 0.549 without it — no difference. A model whose loop settles on zero net turning holds a concentrated heading whether or not there is a cue, exactly as the frozen retina held a bar in Sec. 2.14, so a concentrated heading is not evidence of anemotaxis. Airflow makes the turning steadier without making the animal point anywhere in particular.

Two negative controls bound it. A wind that meanders at 0.25 Hz, introduced on the expectation that a steady signal would be removed by the decoder's baseline, stabilises nothing (-0.019 ± 0.034 spin): holding a course requires a reference that does not move. And removing the baseline subtraction altogether collapses heading concentration from 0.639 to 0.032, the standing DNp15 asymmetry then dominating, so that subtraction is what makes orientation possible rather than what prevents it.

2.16 Where the activity stops

Three circuits have now been asked for a computation and failed, and each was explained separately: LC10a silent so fixation fails, LPLC2 unresponsive so looming fails, EPG silent so the compass fails (Sec. 2.15). One measurement is worth more than three accounts. Over a closed-loop episode in the configuration the behavioural tables use, accumulating every neuron's output per simulation substep rather than per game tic, the distribution is:

super class	cells	mean (Hz)	median	silent
sensory	16,938	0.64	0.00	96.0%
optic	77,873	140.15	182.97	11.4%
visual projection	7,684	155.18	117.68	32.2%
central	32,381	18.22	0.00	85.3%
ascending	1,750	0.52	0.00	98.7%
descending	1,305	49.66	0.00	65.8%

Vision leaves the optic lobe without difficulty. The visual projection neurons are the *most* active population in the model: LC11 at 365 Hz, LC4 at 242, LPLC2 at 207. What fails is the crossing into the central complex, and it fails at an identifiable junction. ER, the ring neurons that carry visual input into the ellipsoid body, fire at 0.00 Hz; Delta7, the inhibition that sharpens a heading bump, at 0.00; EPG at 0.30 and PEN at 2.62. The compass is therefore not a circuit that computes the wrong thing — it is a circuit that is never driven.

This localises the per-cell-type gain question the study raises. Rather than a diffuse claim that unspecified gains hold the model back everywhere, the measurement points at the synapses from visual projection neurons onto ring neurons, and predicts that any behaviour requiring the central complex — menotaxis, path integration, visual place learning — will fail in the same way and for the same reason.

It also corrects a claim made earlier in this work. We have described LPLC2 as silent; that figure comes from the synthetic looming rig, and in closed loop the population fires at 207 Hz. The looming failure is a failure to *discriminate* expansion from a size-matched static disc (Table 3), not an absence of activity. LC10a at 8.77 Hz remains genuinely quiet, some thirty-fold below its neighbours in the same episode, so the fixation account stands as written.

2.17 Two interface errors, and what correcting them changes

Plotting each ommatidial column at the direction it looks, rather than on the lattice, shows a slanted parallelogram slanted the same way in both eyes, where a real right eye mirrors the left. Methods gives the two errors this exposed: an eye map read with the wrong hexagonal convention and mirrored in neither eye, and a steering sign that turned the model away from its more active steering neuron. Every result above was measured through that interface. Both corrections are options rather than defaults, so those results reproduce bit for bit, and the question here is only what the corrections buy.

For the reflex they should matter most to, nothing. In a tethered drum — the commands recorded and discarded, the body turning only as the drum dictates and not walking, so that a counter-turn cannot cancel its own stimulus — ten seeds per speed give no directional response at either level we can read. The commanded yaw modulation, taken against each seed’s own no-drum baseline, is $+0.06 \pm 0.18$, $+0.09 \pm 0.18$, $+0.39 \pm 0.19$ and -0.05 ± 0.18 at 1, 2, 4 and 8 deg/tic; the DNP15 left-right difference that precedes the decoder’s deadzone and baseline filter is $+0.39 \pm 0.50$, $+0.20 \pm 0.52$, -0.14 ± 0.46 and $+0.20 \pm 0.57$ Hz. One cell of sixteen separates, points the wrong way, and does not reverse under a mirrored retina. The frozen-retina arm returns exactly zero, which is the check that the bookkeeping is sound: with an unchanging image both half-cycles are identical.

Closed loop looks more encouraging and is misleading. Flipping the sign chain in exactly one place — a mirrored retina, or the uncorrected steering sign — produces a large speed-tuned modulation ($+1.26 \pm 0.29$ and $+0.68 \pm 0.45$ at 1 deg/tic) in the direction that *adds* to the imposed rotation, while flipping it twice or not at all produces nothing. That is positive feedback amplifying a wrong-way turn: the arm whose turn opposes the drum reduces its own stimulus, the arm that follows increases it, and the asymmetry is a property of the loop rather than of the pathway. It is the reason the tethered version is the one we report.

Rebuilding the arena for the animal moves the response without moving the dissociation. In the fly arena, tethered, the drum modulates DNP15 by $+1.22 \pm 1.11$ Hz at 1 deg/tic and the command by -0.49 ± 0.54 , against $+0.39 \pm 0.50$ and $+0.06 \pm 0.18$ in the sky arena and exactly zero with a frozen retina, so the response is visually driven and roughly three times larger. It is still not directional: the mirrored arm gives the same sign and nearly the same magnitude ($+0.83 \pm 0.99$ Hz), where a genuine optomotor response must reverse.

What the corrected map does buy is visible in the open-loop grating. Direction separation at the horizontal system, rightward minus leftward drift, rises from 1.625 to 3.375 Hz in the configuration the behavioural runs use, and is 0.500 against 0.375 Hz in the frozen model. Per-cell selectivity is unchanged ($|DSI| \approx 0.007$), as it must be: the map fixes where a column looks, not what a detector computes. A doubled population signal of 3 Hz against a standing left–right asymmetry near 20 Hz does not reach behaviour, and the ceiling of Sec. 2.7 is what it runs into.

3 Discussion

The dissociation we observe is not between easy and hard computations, nor between small and large circuits. It is between computations that a linear summation followed by a threshold can express and those it cannot. Gustatory and innate-olfactory valence are, at the level the connectome constrains them, weighted sums with signs: sugar excites a motor pool, bitter inhibits it, an odour drives a valence region that inhibits an escape neuron. Those transfer. Motion direction is a product of a delayed and an undelayed signal, and it does not.

This result is easily over-read as a claim about connectomics, and we state its scope precisely. The connectome contains the correlator: the geometry, the mirrored axes, the phase offset and the travelling wave that drives them are all present and measured, and read at the population stage the assembled circuit does produce a directional signal that ablating its motion detectors removes. What the frozen model fails to supply is the arithmetic that combines them. A product requires a nonlinear interaction between the arms, which in a conductance-based neuron can only come from inhibitory division, and that division is cancelled by the hyperpolarisation it produces (Sec. 2.7). The cell sums where it must multiply. Removing the hyperpolarising component entirely, which is the condition most favourable to a divisive account, abolishes the residual selectivity instead of raising it: what survives is a threshold conjunction rather than a product, which is a poor substitute and measures as one. This accounts for the twenty candidate remedies that changed nothing: each addressed a component that was not defective. It also accounts for the one remedy that appeared promising and was not, the $37\times$ imbalance between the arms, which is a real property of the wiring but not the operative cause, since correcting it by short-term depression and by transplanting an optimized model’s gains both lowered selectivity.

Restoring the interaction requires raising the inhibitory-to-excitatory conductance ratio, and the constraint that emerges is the paper’s substantive claim. No single value serves the whole brain. The optic lobe requires a ratio that inverts the gustatory circuit, because 51% of inhibitory synapses terminate on inhibitory neurons and uniform scaling amplifies disinhibition equally. The same factor applied regionally leaves the gustatory circuit numerically unchanged. A single global synaptic gain is therefore not an approximation to be refined but an assumption that cannot hold, and the connectome is silent on the point: synapse counts constrain relative weights within a target, not the conductance scale across regions.

Placed beside Lappalainen et al. [2024] this is informative rather than merely negative. That model shares our wiring and recovers motion selectivity by optimizing its unknown parameters, where ours fixes them to a single published value. The difference between the outcomes measures what optimization supplies. Transplanting its per-type gains directly does not transfer the benefit, so what it supplies is not gain as a lookup table; the regional result suggests it is at least in part the freedom for conductance scale to differ between circuits, which per-type parameters express implicitly. Published physiology is consistent with this reading, reporting GABAergic conductances around $2.5\times$ cholinergic and inhibitory-to-excitatory ratios of 7.7 and 10.9 in adult lobula plate neurons [Su and O’Dowd, 2003, Groschner et al., 2022]. Applying those measured ratios uniformly, however, is not sufficient: it improves selectivity and breaks olfaction (Table 5), which indicates that correct per-receptor values and regional assignment are separate requirements.

The behavioural measurements do not extend the result, and we take that seriously rather than reporting the subset that would. On thirty held-out seeds no parameterisation separates from its command-matched control by more than the frozen model does, and the regionally scaled model matches the frozen model at six of eighteen measures. Restoring the arithmetic that direction selectivity requires is therefore not sufficient to produce the behaviour that selectivity supports, at least not measurably through this environment.

We initially read this as ambiguous between a manipulation too small and a measurement too coarse. The population readout resolves it, and not in favour of either. At the optomotor output stage the directional signal varies by an order of magnitude and by sign across these configurations, from +24.1 Hz to −1.5; behaviour does not follow, and the same two configurations return four and six of eighteen. The manipulation was not too small at that site, and the quantity that fails to appear is not selectivity but its transmission to the readout. Two mechanically unrelated interventions, the regional scale and the transfer function, produce one behavioural fingerprint in metrics of locomotion with the vision–steering correlation flat, which is what altered motor statistics look like rather than sensory control.

The measurement is nonetheless implicated. Health collected varies by roughly 13 units across seeds against an effect of a few units, the power analysis puts the required sample near $n \approx 106$, and two runs of one configuration returned 3.2 and 11.2. An instrument that cannot resolve a population signal differing by an order of magnitude between conditions is not the instrument for this question.

The readout supplies the rest of the answer, and it revises what the negative result is about. DNa02 receives 2.3% of its input from visual populations and 0.34% from the horizontal system, and the pathway the optomotor literature describes, from the horizontal system to DNa02, carries 75 of 7,297 units of signed weight in this reconstruction. The signal is not failing to be computed, nor failing to leave the optic lobe: it arrives at DNp15, the horizontal system’s actual descending target, at 32 Hz. What fails is the correspondence between the cell a decoder was pointed at, chosen from the literature, and the cell this wiring diagram routes the computation to.

The published characterisations settle which cell should have been read. DNa02 steers goal-directed walking under central-complex control [Rayshubskiy et al., 2025]; DNp15 is the optomotor descending neuron driven by the horizontal system [Suver et al., 2016, Namiki et al., 2018]. Decoding the latter, with everything else fixed, quadruples the coupling between steering and retinal luminance, and mirroring the retina abolishes that gain rather than reversing it. The coupling is vision-dependent without the sign structure of an optomotor response, and task performance is unchanged. The more useful statement is methodological: in a frozen model the readout is a parameter like any other, and it was the least examined one here, chosen from the literature for a different pathway than the stimulus being delivered. The same applies at the population level, where 79.8% of descending neurons never fire and the decoded pair runs near saturation, so the motor stage constrains behavioural claims independently of anything the sensory side computes. What can be said is that episode-level behavioural measures in a stochastic environment are a weak instrument for this question, and that a closed-loop connectome study should establish their power before resting a claim on them.

Two further readings should be excluded. First, none of this is a claim about whether a fly brain can play a video game. The environment lies outside the animal’s evolutionary experience and the model has no ventral nerve cord, so the behavioural measures were intended as an assay of whether a sensorimotor transformation is present at all, never as a measure of competence, and the result above is that they are not sensitive enough to serve even that purpose. Second, the olfaction result is not evidence of enemy detection: the odour magnitude is supplied by the engine. What both results measure is transfer: whether a transformation implied by the wiring appears in a running system.

3.1 Limitations

- **No ventral nerve cord.** FAFB is brain-only, so the game’s movement code substitutes for legs. All behavioural *timescales* are the engine’s; claims are restricted to descending-neuron activity.
- **No gap junctions.** EM reconstruction captures chemical synapses; the giant fibre’s fastest output is electrical and therefore invisible here.
- **No neuromodulation.** Dopamine, octopamine and serotonin are flattened to fast excitation.
- **$n = 1$ brain, and its halves are not mirror images.** DNa02 sits $\sim 2\times$ higher on the right than the left, which is a standing turn command until removed by a slow baseline.
- **Angular scale is calibrated, not measured.** The lattice is real but its axes are not isotropic in visual angle, so each is scaled to the published field of view. Absolute angular claims inherit this.

- **The gustatory control is partly true by construction**, since W_{syn} is fitted on it.
- **The agent walked backwards.** The walk command is BPN—MDN, whose standing value was a forward walk when the decoder was designed because MDN was silent. The optic gain of 16 that the behavioural configuration requires drives MDN from 0 to 107 Hz, so in every behavioural result reported here the command was negative on 97% of tics and the agent spent 37–68% of its time reversed against a wall. The matched control, sharing its command statistics, reversed too, so the comparisons are sound; the behaviour is not fly-like. Reading MDN as bursts above its own running level, as a fly’s is [Bidaye et al., 2014], and counting antennal contact only on forward pushes restores forward walking (0% reversed) and roughly doubles exploration over three levels. Re-run with it, the behavioural advantage survives only in the fly arena (Sec. 2.13).
- **The visual interface carried two errors** (Sec. 2.17): an eye map read with the wrong hexagonal convention and unmirrored between the eyes, and a reversed steering sign. Every result reported here was measured through it. Correcting both doubles the population direction signal, leaves per-cell selectivity unchanged, and produces no optomotor response, so the conclusions are unchanged; but any result that depends on the sign structure of binocular optic flow should be read with this in mind.
- **One rendering engine.** All visual results come from Doom, and all but the full-eye runs from one viewport geometry; the full-eye runs add a second arena and field of view, but we have not tested whether a different visual statistic changes the arm imbalance. Doom’s sky is not drawn in true perspective, so the three views of the full eye disagree slightly above the horizon.

4 Methods

4.1 Connectome data

The connectome is FlyWire FAFB v783 [Dorkenwald et al., 2024, Schlegel et al., 2024], frozen October 2023. Connectivity from the Buhmann synapse pipeline [Buhmann et al., 2021] with transmitter predictions from Eckstein et al. [2024]. Edges thresholded at five synapses. Cell types from the consolidated typing table; BPN and aIPg exist only in the free-text community label file. The data are CC BY-NC 4.0 and are not redistributed with our code.

4.2 Neuron model

Neurons are leaky integrate-and-fire with the parameters listed in Sec. 2.1. The subthreshold system $\tau \dot{v} = (V_{\text{rest}} - v) + g$, with g decaying at τ_{syn} , is integrated exactly:

$$v(t+dt) = v_0 A + g_0 B + g_{\text{ext}}(1 - A), \quad A = e^{-dt/\tau_{\text{mem}}}, \quad B = \frac{\tau_{\text{syn}}}{\tau_{\text{syn}} - \tau_{\text{mem}}} (e^{-dt/\tau_{\text{syn}}} - A),$$

matching the scheme used by Brian 2 [Stimberg et al., 2019], against which we validate to 0.03%.

4.3 Departures from the reference model

Six components depart from the reference model of Shiu et al. [2024]. Each is switchable and each carries a measured justification. (i) *Histamine override* on photoreceptor output, as above. (ii) *Conductance-based synapses*: $\tau \dot{v} = (V_{\text{rest}} - v) + g_e(E_e - v) + g_i(E_i - v) + I$ with $E_e = 0$, $E_i = -70$ mV, verified divisive (shunting reduces input–output slope by $> 40\%$ where subtraction preserves it). (iii) *Per-type conduction delays* via a multi-tap ring buffer, since a single global delay cannot express a correlator. (iv) *Graded non-spiking units* for all 77,873 optic neurons, output rectified-linear in $(v - V_{\text{rest}})$ capped at 200 Hz. (v) *Weber contrast adaptation* per column. (vi) L3 kept sustained while L1/L2 adapt: from a stationary agent the rendered scene changes by 0.0003 in mean luminance between frames, so a fully adapting retina washes out and the loop cannot start.

4.4 Visual input and the environment

The environment is modified where its rendering conflicts with fly photoreception. Modifying the game is not a modification of the model, and three properties of Doom are wrong for this animal in ways that suppress the signal under test. *Gamma*: the engine writes gamma-encoded sRGB, but

photoreceptors integrate linear flux and the acceptance function is a spatial integral over it, so the Gaussian pre-blur must run in linear light; undoing the encoding first raises retinal contrast by 37%. *Spectrum*: R1-6 carries Rh1, which is blue-green sensitive and nearly blind at Doom’s red primary, while Doom is painted in browns: weighting the frame by R1-6 sensitivity therefore *lowers* retinal signal relative to human luma. *Palette*: because R1-6 are a single spectral class the motion pathway is monochromatic, so hue is not a quantity it can use, and we re-emit the scene’s achromatic structure in blue-green: equivalent to repainting every surface, and the same move a fly-vision laboratory makes when it builds an arena from green emitters. That recovers the lost signal and then some, raising retinal standard deviation from 0.020 to 0.033, measured open loop on 100 identical frames.

Full visual field. Each eye’s lattice spans $170 \times 150^\circ$ and the eyes point 40° outward, about 250° together, while one planar view cannot exceed 180° . For the full-eye runs we render three 170° views at 1280×1024 , forward and 90° to either side, and each column reads the view nearest its own direction. The engine draws one view per tic and ties the camera to the heading, so each side view comes from a separate engine instance that loads the main game’s saved state every tic and advances one tic with the agent’s action rotated by the offset; it renders the same world, not a replica. Three engine behaviours had to be handled: a console save lands on alternate tics and holds the state one tic before the tic that wrote it, so the intervening actions are replayed; loading itself runs a tic with the last-set action, which is therefore set first; and movement commands are floored to integers, so the quarter turn is applied as an exact integer permutation. At zero offset the side view is pixel-identical to the main view over whole episodes; at $\pm 90^\circ$, columns in the overlap read the same scene from both views ($r = 0.97$ and 0.99). Two corrections accompany the wider field. A planar view places elevation e at $\tan e / \cos a$, not $\tan e$; the earlier mapping, used for every result except the full-eye runs, points a column at an elevation off by a median 1.8° (8.0° beyond 40° of azimuth) and counts 148 columns as inside the viewport that are not. And since a 170° view magnifies its periphery up to 130-fold, each column takes its acceptance blur from the level of an image pyramid matched to its local scale. The engine’s head-bob is pinned at its default, because ViZDoom persists console variables across sessions and a diagnostic had silently set it to zero.

Where each column looks. The map above places a column by stretching the lattice’s own Cartesian coordinates onto the published field of view. That map, used for every result except where noted, contains two errors, both measured rather than argued. First, FlyWire’s (p, q) are axial hexagonal coordinates whose axes are 120° apart, not 60° : read at 60 each eye is a 51×17 sliver, read at 120 it is 32×27 , and the four T4/T5 preferred directions come out 95° apart rather than 110° . Second, both hemispheres carry the same anatomical convention, so one eye’s map must be mirrored, and neither was: for a rotating world both eyes then report the same local direction, which is the one stimulus an optomotor response must distinguish, while forward translation reports opposite ones.

The corrected map takes its axes from the wiring. For each subtype and hemisphere we take the synapse-weighted offset between the columns supplying the tip of a dendrite and those supplying its base (Mi9 against Mi4/C3 for T4; Tm9 against Tm1/Tm2 for T5), which estimates the preferred direction from anatomy alone. T4 and T5 agree within $1\text{--}13^\circ$ on all four directions in both eyes, placing front-to-back along $-h$ and upward along $+v$, where $h = (q - p)\sqrt{3}/2$ and $v = (p + q)/2$. The R7/R8 columns presynaptic to the dorsal-rim DmDRA cells lie on the $+v$ side, which fixes “up” without reference to motion. Both agree with the published eye map, in which $+h$ “points towards the posterior medulla, corresponding to anterior on the eye because of the optic chiasm” [Zhao et al., 2025]. Each eye is then placed as measured: from 10° into the opposite hemifield in front to 155° behind, and from 70° below the horizon to directly above [Zhao et al., 2025, Kempainen et al., 2022], which supersedes the $\pm 40^\circ$ splay. 98% of columns then fall inside the three views; the remainder look above 82° , which no planar projection renders.

Sky arena. The standard arena has a dark stone ceiling, so the dorsal half of every eye viewed near-black. The full-eye runs use the same map with the ceiling replaced by open sky at the same height and a generated daylight texture, which raises the mean brightness of the top quarter of the view from 10.8 to 196 (of 255).

Fly arena. A further variant addresses the rest of the field on the same argument. Doom’s wall detail sits near one map unit, far below the 5.4° acceptance angle, so the eye averages it to flat grey and the

only wall contrast left is the silhouette against the sky; its floor is a dark liquid flat at 0.023 luminance against the sky's 0.71, with a per-lens standard deviation of 0.008, which gives the ventral field nothing to track. The replacements are a vertical square-wave grating of period 128 map units, which spans 18–37° at the distances walls sit at and so brackets the 20–30° optimum for the optomotor response, and a bright irregular ground with features of 32–64 units; sector lighting is set flat and maximum, since Doom shades sectors to fake a depth cue for an eye that infers depth from parallax. Frame-to-frame change in per-lens luminance, which is the quantity a correlator integrates, rises 1.9× over the whole eye and 3.5× in the ventral field, while static contrast falls slightly: what it loses is a large static edge that no motion detector can use.

Render resolution. Each column integrates a 6.3° acceptance function, so most of a 1280 × 1024 frame is averaged away before the retina sees it. At 640 × 480 per-column luminance correlates 0.995 with the full render (mean absolute difference 0.014 against a 0.245 spread across the eye) and the same 4,540 of 4,541 columns fall inside the views, at 1.7 times the speed.

None of it releases LPLC2. We note the engine's own display controls are not a route either: the engine's gamma, contrast, brightness and saturation controls all leave the rendered buffer bit-identical, because ViZDoom exposes it upstream of display post-processing, and `r_visibility`, which does take effect, moves frame mean brightness without changing the distance-to-brightness relation by more than 3–6% over a 40-tic approach. Every one of these applies a common scale to both arms, so none of them alters the ratio between them, and none addresses the cancellation identified in Sec. 2.7.

4.5 Rate readout and motor decoding

Firing rate is read out by an exponential filter over per-step output. Because that output is rate × dt , recovering Hz requires division by dt ; dividing by τ instead scales every reported rate by 160× and is a failure mode we hit and document. A 12-tic warmup precedes baseline estimation, and a $\tau = 3$ s baseline removes the standing left–right asymmetry.

The steering sign was wrong, and every result except where noted inherits it. The decoder takes the left-minus-right difference of the steering pair and sends it to `TURN_LEFT_RIGHT_DELTA`, on the assumption that the engine treats positive as a left turn. Measured, a +10 command lowers the heading by 10° and displaces the rendered scene 28 px to the left at 90° over 320 px, which is a right turn. The model therefore steered *away* from the side whose steering neuron fired harder, where both DNa02 [Rayshubskiy et al., 2025] and DNp15 steer toward their own side. The corrected sign is available as an option and used only where stated.

4.6 Statistical analysis

The olfaction comparison uses 10 independent environment seeds. Each seed contributes one paired on/off difference, averaged over 320 tics after a 40-tic settling period that lets the rate filter converge. We report the mean and standard deviation across seeds together with a two-sided sign test on the number agreeing with the mean direction ($p = 2 \sum_{i=k}^n \binom{n}{i} / 2^n$; $p = 0.002$ at 10/10, $p = 0.021$ at 9/10). We prefer the sign test to a t -test because per-seed effect magnitude depends on how closely enemies happen to approach, which is neither normal nor identically distributed across seeds: it is a property of the episode, not of the model. Note that sign consistency alone is a weak claim when an effect is small, which is why we read the LC4 control on magnitude rather than on its p -value. The degree-preserving shuffle permutes postsynaptic targets while preserving each neuron's in-degree, and is run through the identical protocol with the identical seeds. Table and figure are generated from the same serialized run data, so they cannot diverge.

The environment is not bit-reproducible across processes, so the magnitudes carry run-to-run variation of a few percent and the sign counts are less stable than a single execution suggests. We can put a number on the latter. The odour-source table contained an error: `CustomMedikit`, the pickup actor used by one scenario, was absent from the food list, and unlisted actors fell through to the pheromone channel, so inanimate scenery was being smelled as conspecifics. Correcting it left the lateral horn response essentially unchanged in size but moved DNp01 and DNa02 from 10/10 agreeing seeds to 9/10. Sign counts near ceiling should therefore be read as p values with one seed of slack, not as invariants.

Power of episode-level behavioural measures. Behaviour in this environment is far noisier than the open-loop rate measurements, and we flag the consequence because it constrains what any closed-loop connectome study can claim. Health collected per episode has a standard deviation of roughly 13 units across environment seeds, so at $n = 5$ only differences of about 20 units are resolvable; two executions of one identical configuration returned means of 3.2 and 11.2. Detecting a 5-unit difference at 80% power would require $n \approx 106$. The grid of Sec. 2.13 shows this is not a conservative bound but a binding one: a population signal differing by an order of magnitude and by sign between conditions produced no separation in these measures. We therefore report no behavioural effect from fewer than 30 seeds, quote 95% confidence intervals rather than standard deviations for these measures, and treat the interval rather than the point estimate as the result. Rate-level measurements, which are paired and open loop, are not subject to this and are reported as above.

4.7 Negative results and deferred work

Recording what was tried and failed is as much a part of the result as what worked, so we list both.

Tested, and does not rescue visual motion: per-type conduction delays (1.8–240 ms); per-neuron membrane time constants (5–200 ms); conductance-based shunting synapses, which are now the default rather than a candidate; a forced high-conductance state; population pooling; graded non-spiking optic-lobe units; tonic optic-lobe bias from 0 to 7.5 mV; photoreceptor drive, graded rate ceiling and Weber contrast gain; injection at the photoreceptors rather than the lamina; field of view from 110° to 170°; sRGB linearisation; R1-6 spectral weighting; repainting the environment into the fly’s band; and the engine’s own rendering controls. Most of these share a structure: each applies a scale to both arms of a ratio, and such a scale cancels. That account is incomplete on its own, since the *asymmetric* interventions listed below do not cancel and failed as well; the common cause is given in Sec. 2.7.

Corrected during this work, all on the input side: the retina mapped azimuth linearly onto a perspective projection, misplacing a column at 10° of gaze by 87%; the field of view was never applied, because ViZDoom ignores it as a command-line argument, so the model ran at the engine default rather than the reported value; frames were held across all 57 substeps of a tic, presenting a correlator with an impulse train aliased against its own 80 ms delay line; and, consequently, Weber adaptation advanced once per tic rather than once per substep, giving it an effective time constant of 14 s in place of 0.25 s. Each of these corrupts precisely the cue under test. None of them changed the outcome, which is itself informative: the negative result does not rest on a badly delivered stimulus.

Two silences and a control defect, diagnosed rather than listed. LC10a fires exactly zero across the fixation test, which reports a left-right difference of 0.000; its input is inhibition-dominated at $g_i/g_e = 1.85$, which accounts for it. LPLC2 is also silent at 0.75, beside LC4 at 0.59 which fires, so that ratio does not account for this one and we do not claim it does. In the looming test the static control was the expansion’s final radius held for the whole window, 70° against a sweep beginning at 5°, so it delivered more total drive than the stimulus it controlled for, and LC4 responding more to static than to looming was a defect of that comparison rather than of adaptation. That is now repaired: the acceptance check is scored against a static disc of the sweep’s *mean* angular size, which delivers the same total drive, and against it LC4 gives 26.42 against 26.44 Hz (Table 3). The verdict is unchanged and no longer rests on a comparison that could not have passed.

Bounding the accessible space: the inhibitory reversal potential from −52 to −100 mV, which peaks at the published value and collapses when inhibition is made purely divisive; the regional inhibitory scale from 1 to 3.5; published per-receptor ratios both globally and regionally; and the optic gain from 4 to 64, monotone but shallow and at rates well above physiological. Largest selectivity found anywhere: −0.0792.

Also tested, and also negative: presynaptic short-term depression as an arm rebalancer; transplantation of 571 per-type gains from an optimized model; ablation of CT1 (12,601 edges); ablation of one inhibitory flank; a uniform optic-lobe gain to $\times 32$ and the same gain targeted onto T4/T5 to $\times 20$, both of which multiply excitatory and inhibitory synapses alike and are therefore distinct from the inhibitory-to-excitatory ratio of Sec. 2.9; matching the network’s firing-rate regime to that of the isolated correlator; a supralinear, voltage-dependent excitatory conductance on the dendrite, which does nothing a flat gain of the same mean size does not (−0.0425 against −0.0428); and replacing the slow arm’s fixed

conduction delay with a one-pole low-pass filter, which is worse at every time constant tried and best where its time constant equals the delay it replaces (0.039 against 0.064). *Positive control added:* a hand-built two-input correlator in the same simulator reaches $DSI = -0.79 \pm 0.017$, which is what licenses us to call the network result a failure of the wired model rather than of the method. Rebuilt independently, the same construction reaches -0.26 at one configuration and -1.00 once the real fan-in and spatial spread are added; we quote the original figure but note that the magnitude depends on the operating point chosen, and that the qualitative claim (that this model class computes direction) is the part the argument rests on.

Two analysis pitfalls, recorded for reproduction. Both yield plausible, self-consistent numbers and neither announces itself. The first is a configuration option applied in one code path but not in others that construct their own copy of the same mask: measurements taken through those paths characterise the default model rather than the requested one. We now set the option in a single place used by every path. The second is summarising a cell's input geometry by pooling Mi9 with Mi4. These sit on *opposite* dendritic flanks, so their centroid inverts with their relative weight; pooling them manufactures an apparent bimodality in the axis distribution and an apparent $97\text{--}114^\circ$ mirror-pair defect where none exists. Taking Mi9 alone: the axes are coherent ($R = 0.67\text{--}0.93$) and the mirror pairs are $182\text{--}194^\circ$ apart, as they should be. T4's input geometry is not defective.

Deferred. Per-cell-type synaptic gains remain unfitted, because fitting them is the independent variable of this study rather than a repair to it, a position unchanged by the transplant result above, which tests transfer rather than fitting. We note the target is not a single set of numbers: the optimized model of Lappalainen et al. [2024] carries 734 free parameters, 604 of them one gain per type-to-type connection, and its authors report an ensemble of 50 solutions of comparable task performance, so connectome plus task does not identify them uniquely. Also deferred: gap junctions, absent from the reconstruction and carrying the giant fibre's fastest output; neuromodulation, flattened here to fast transmission; and a ventral nerve cord, without which every behavioural timescale is the engine's.

Data availability

The connectome data are FlyWire FAFB v783 [Dorkenwald et al., 2024, Schlegel et al., 2024], licensed CC BY-NC 4.0. They are not redistributed with this work and must be obtained from FlyWire directly; a fetch script is provided with the code. The per-cell-type gains used in the transplant ablation are derived from the published flyvis ensemble of Lappalainen et al. [2024]. All serialized run data behind the tables and figures are included in the code repository, and every table and figure is generated from those files rather than transcribed, so the two cannot diverge.

Code availability

Code is available at <https://github.com/mutkuoz/flydoom> under the MIT licence, including the simulator, all ten experiments, and the scripts that generate every table and figure. The environment uses ViZDoom [Kempka et al., 2016] with the freely licensed Freedoom asset set, so no commercial game data are required.

DOOM is a trademark of id Software LLC / ZeniMax Media Inc. This work is unaffiliated with and unendorsed by them, and the name is used here only to identify the software environment.

Declaration of generative AI and AI-assisted technologies

During the preparation of this work, the author utilized large language models, specifically Claude Opus 5 by Anthropic, as an assistive technology under direct instruction. The model was employed to execute code generation tasks in Python against author-defined specifications, and to draft and refine the language and structure of the manuscript. All AI-assisted output, both code and text, was reviewed, tested and edited by the author. No quantity reported here was produced by a language model: every number comes from executing the released code, and every table and figure is emitted from the serialized run data by a script. The author retains sole responsibility for the validity, integrity and final content of this publication.

Acknowledgements

This work uses the FlyWire FAFB connectome, produced by the FlyWire Consortium and the many proofreaders whose annotations make cell-type resolution possible.

References

- Salil S. Bidaye, Christian Machacek, Yang Wu, and Barry J. Dickson. Neuronal control of *Drosophila* walking direction. *Science*, 344(6179):97–101, 2014. doi: 10.1126/science.1249964.
- Salil S. Bidaye, Meghan Laturney, Amy K. Chang, Yuejiang Liu, Till Bockemühl, Ansgar Büschges, and Kristin Scott. Two brain pathways initiate distinct forward walking programs in *Drosophila*. *Neuron*, 108(3):469–485.e8, 2020. doi: 10.1016/j.neuron.2020.07.032.
- Alexander Borst, Juergen Haag, and Dierk F. Reiff. Fly motion vision. *Annual Review of Neuroscience*, 33(1):49–70, 2010. doi: 10.1146/annurev-neuro-060909-153155.
- Julia Buhmann, Arlo Sheridan, Caroline Malin-Mayor, Philipp Schlegel, Stephan Gerhard, Tom Kazimiers, Renate Krause, Tri M. Nguyen, et al. Automatic detection of synaptic partners in a whole-brain *Drosophila* electron microscopy data set. *Nature Methods*, 18(7):771–774, 2021. doi: 10.1038/s41592-021-01183-7.
- Thomas A. Cleland. Inhibitory glutamate receptor channels. *Molecular Neurobiology*, 13(2):97–136, 1996. doi: 10.1007/bf02740637.
- Michael-John Dolan, Shahar Frechter, Alexander Shakeel Bates, Chuntao Dan, Paavo Huoviala, Ruairí JV Roberts, Philipp Schlegel, Serene Dhawan, et al. Neurogenetic dissection of the *Drosophila* lateral horn reveals major outputs, diverse behavioural functions, and interactions with the mushroom body. *eLife*, 8:e43079, 2019. doi: 10.7554/elife.43079.
- Sven Dorkenwald, Arie Matsliah, Amy R. Sterling, Philipp Schlegel, Szi-chieh Yu, Claire E. McKellar, Albert Lin, Marta Costa, et al. Neuronal wiring diagram of an adult brain. *Nature*, 634(8032):124–138, 2024. doi: 10.1038/s41586-024-07558-y.
- Nils Eckstein, Alexander Shakeel Bates, Andrew Champion, Michelle Du, Yijie Yin, Philipp Schlegel, Alicia Kun-Yang Lu, Thomson Rymer, et al. Neurotransmitter classification from electron microscopy images at synaptic sites in *Drosophila melanogaster*. *Cell*, 187(10):2574–2594.e23, 2024. doi: 10.1016/j.cell.2024.03.016.
- Lukas N. Groschner, Jonatan G. Malis, Birte Zuidinga, and Alexander Borst. A biophysical account of multiplication by a single neuron. *Nature*, 603(7899):119–123, 2022.
- Eyal Gruntman, Sandro Romani, and Michael B. Reiser. Simple integration of fast excitation and offset, delayed inhibition computes directional selectivity in *Drosophila*. *Nature Neuroscience*, 21(2):250–257, 2018. doi: 10.1038/s41593-017-0046-4.
- Juergen Haag, Alexander Arenz, Etienne Serbe, Fabrizio Gabbiani, and Alexander Borst. Complementary mechanisms create direction selectivity in the fly. *eLife*, 5:e17421, 2016. doi: 10.7554/elife.17421.
- R. C. Hardie. A histamine-activated chloride channel involved in neurotransmission at a photoreceptor synapse. *Nature*, 339(6227):704–706, 1989. doi: 10.1038/339704a0.
- B. Hassenstein and W. Reichardt. Systemtheoretische Analyse der Zeit-, Reihenfolgen- und Vorzeichenauswertung bei der Bewegungspertzeption des Rüsselkäfers *Chlorophanus*. *Zeitschrift für Naturforschung B*, 11(9-10):513–524, 1956. doi: 10.1515/znb-1956-9-1004.
- Gregory S.X.E. Jefferis, Christopher J. Potter, Alexander M. Chan, Elizabeth C. Marin, Torsten Rohlfsing, Calvin R. Maurer, and Liqun Luo. Comprehensive maps of *Drosophila* higher olfactory centers: Spatially segregated fruit and pheromone representation. *Cell*, 128(6):1187–1203, 2007. doi: 10.1016/j.cell.2007.01.040.

- Kyobi S. Kakaria and Benjamin L. de Bivort. Ring attractor dynamics emerge from a spiking model of the entire protocerebral bridge. *Frontiers in Behavioral Neuroscience*, 11:8, 2017. doi: 10.3389/fnbeh.2017.00008.
- Michal Kempka, Marek Wydmuch, Grzegorz Runc, Jakub Toczek, and Wojciech Jaskowski. ViZDoom: A Doom-based AI research platform for visual reinforcement learning. In *2016 IEEE Conference on Computational Intelligence and Games (CIG)*, pages 1–8, 2016. doi: 10.1109/cig.2016.7860433.
- Joni Kemppainen, Ben Scales, Keivan Razban Haghighi, Jouni Takalo, Neveen Mansour, James McManus, Gabor Leko, Paulus Saari, James Hurcomb, Andra Antohi, Jussi-Petteri Suuronen, Florence Blanchard, Alexey A. Polilov, Tuomas Takalo, Mark Hensman, Roger C. Hardie, and Mikko Juusola. Binocular mirror-symmetric microsaccadic sampling enables *Drosophila* hyperacute 3D vision. *Proceedings of the National Academy of Sciences*, 119(12):e2109717119, 2022. doi: 10.1073/pnas.2109717119.
- Amina Kurtovic, Alexandre Widmer, and Barry J. Dickson. A single class of olfactory neurons mediates behavioural responses to a *Drosophila* sex pheromone. *Nature*, 446(7135):542–546, 2007. doi: 10.1038/nature05672.
- Janne K. Lappalainen, Fabian D. Tschopp, Sridhama Prakhya, Mason McGill, Aljoscha Nern, Kazunori Shinomiya, Shin-ya Takemura, Eyal Gruntman, et al. Connectome-constrained networks predict neural activity across the fly visual system. *Nature*, 634(8036):1132–1140, 2024. doi: 10.1038/s41586-024-07939-3.
- Matthew S. Maisak, Juergen Haag, Georg Ammer, Etienne Serbe, Matthias Meier, Aljoscha Leonhardt, Tabea Schilling, Armin Bahl, et al. A directional tuning map of *Drosophila* elementary motion detectors. *Nature*, 500(7461):212–216, 2013. doi: 10.1038/nature12320.
- Shigehiro Namiki, Michael H Dickinson, Allan M Wong, Wyatt Korff, and Gwyneth M Card. The functional organization of descending sensory-motor pathways in *Drosophila*. *eLife*, 7:e34272, 2018. doi: 10.7554/eLife.34272.
- Aleksandr Rayshubskiy, Stephen L Holtz, Alexander S Bates, Quinn X Vanderbeck, Laia Seratosa Capdevila, Victoria Rockwell, and Rachel Wilson. Neural circuit mechanisms for steering control in walking *Drosophila*. *eLife*, 13:RP102230, 2025. doi: 10.7554/eLife.102230.
- Philipp Schlegel, Yijie Yin, Alexander S. Bates, Sven Dorkenwald, Katharina Eichler, Paul Brooks, Daniel S. Han, Marina Gkantia, et al. Whole-brain annotation and multi-connectome cell typing of *Drosophila*. *Nature*, 634(8032):139–152, 2024. doi: 10.1038/s41586-024-07686-5.
- Julia L. Semmelhack and Jing W. Wang. Select *Drosophila* glomeruli mediate innate olfactory attraction and aversion. *Nature*, 459(7244):218–223, 2009. doi: 10.1038/nature07983.
- Philip K. Shiu, Gabriella R. Sterne, Nico Spiller, Romain Franconville, Andrea Sandoval, Joie Zhou, Neha Simha, Chan Hyuk Kang, et al. A *Drosophila* computational brain model reveals sensorimotor processing. *Nature*, 634(8032):210–219, 2024. doi: 10.1038/s41586-024-07763-9.
- Marcel Stimberg, Romain Brette, and Dan FM Goodman. Brian 2, an intuitive and efficient neural simulator. *eLife*, 8:e47314, 2019. doi: 10.7554/elife.47314.
- Hailing Su and Diane K. O’Dowd. Fast synaptic currents in *Drosophila* mushroom body Kenyon cells are mediated by α -bungarotoxin-sensitive nicotinic acetylcholine receptors and picrotoxin-sensitive GABA receptors. *Journal of Neuroscience*, 23(27):9246–9253, 2003.
- Marie P Suver, Ainul Huda, Nastasya Iwasaki, Steven Safarik, and Michael H Dickinson. An array of descending visual interneurons encoding self-motion in *Drosophila*. *Journal of Neuroscience*, 36(46):11768–11780, 2016. doi: 10.1523/JNEUROSCI.2277-16.2016.
- Catherine R von Reyn, Patrick Breads, Martin Y Peek, Grace Zhiyu Zheng, W Ryan Williamson, Alyson L Yee, Anthony Leonardo, and Gwyneth M Card. A spike-timing mechanism for action selection. *Nature Neuroscience*, 17(7):962–970, 2014. doi: 10.1038/nn.3741.

1469 Sibor Wang-Chen, Victor Alfred Stimpfling, Thomas Ka Chung Lam, Pembe Gizem Özdil, Louise
1470 Genoud, Femke Hurtak, and Pavan Ramdya. NeuroMechFly v2: simulating embodied sensorimotor
1471 control in adult *Drosophila*. *Nature Methods*, 21(12):2353–2362, 2024. doi: 10.1038/s41592-024-02497-y.

1472 Arthur Zhao, Eyal Gruntman, Aljoscha Nern, Nirmala A. Iyer, Edward M. Rogers, Sanna Koskela, Igor
1473 Siwanowicz, Marisa Dreher, Miriam A. Flynn, Connor W. Laughland, Henrique D. F. Ludwig, Alex G.
1474 Thomson, Cullen P. Moran, Bruck Sohn, Emily J. McLaughlin, Gerald M. Rubin, and Michael B.
1475 Reiser. Eye structure shapes neuron function in *Drosophila* motion vision. *Nature*, 2025. doi:
1476 10.1038/s41586-025-09276-5.